

# CS 453/698: Software and Systems Security

## **Module: An In-depth Study of Memory Errors**

Lecture: Definition and exploits


Meng Xu (*University of Waterloo*)

Winter 2025

# Outline

- 1 Why study memory errors?
- 2 Demonstration of memory error exploits
- 3 A relatively formal definition of memory errors
- 4 Case study: Heartbleed vulnerability

# Memory errors are prevalent



Source: [BlackHat IL 2019 talk](#) by Matt Miller from Microsoft

Around 70% of all the vulnerabilities in Microsoft products addressed through a security update each year (2006 - 2018) are memory safety issues

# Memory errors are prevalent

High+, impacting stable

Security-related assert

7.1%

Other


23.9%

Other memory unsafety

32.9%

Use-after-free

36.1%




Source: [Chromium Memory Safety Report](#) from Google.

Analysis based on 912 high or critical severity security bugs in Chromium reported in 2015 - 2020

# Memory errors are prevalent

## Memory Safety Vulnerabilities are Disproportionately Severe



Source: Blog post [Memory Safe Languages in Android 13](#) from Google.

Memory safety vulnerabilities disproportionately represent Android's most severe vulnerabilities

# Memory errors can lead to severe consequences



Heartbleed Vulnerability  
(CVE-2014-0610)

# Memory errors can lead to severe consequences



Heartbleed Vulnerability  
(CVE-2014-0610)

- A security bug in version 1.0.1 of OpenSSL, which is a widely used implementation of the Transport Layer Security (TLS) protocol
- It was introduced into OpenSSL in 2012 and publicly disclosed in April 2014
- At the time of disclosure, some 17% (around half a million) of the Internet's secure web servers certified by trusted authorities were believed to be vulnerable to the attack

# Memory errors can lead to severe consequences



Heartbleed Vulnerability  
(CVE-2014-0610)

- The Canada Revenue Agency (CRA) reported a theft of social insurance numbers belonging to 900 taxpayers, and said that they were accessed through an exploit of the bug during a 6-hour period on 8 April 2014.
- After the discovery of the attack, the agency shut down its website and extended the taxpayer filing deadline from 30 April to 5 May.



# Memory errors can lead to severe consequences



Heartbleed Vulnerability  
(CVE-2014-0610)

- The Canada Revenue Agency (CRA) reported a theft of social insurance numbers belonging to 900 taxpayers, and said that they were accessed through an exploit of the bug during a 6-hour period on 8 April 2014.
- After the discovery of the attack, the agency shut down its website and extended the taxpayer filing deadline from 30 April to 5 May.
- On 16 April, the RCMP announced they had charged [a computer science student](#) in relation to the theft with unauthorized use of a computer and mischief in relation to data.

# Heartbleed explanation

SERVER, ARE YOU STILL THERE?  
IF SO, REPLY "POTATO" (6 LETTERS).



This page about "books". User Alice requests  
secure connection using key "4538538374224".  
User Meg wants these 6 letters: POTATO. User  
Ada wants pages about "irl games". Unlocking  
secure records with master key 5130985733435  
leak! (When we leak this message: "P



POTATO



This page about "books". User Alice requests  
secure connection using key "4538538374224".  
User Meg wants these 6 letters: **POTATO**. User  
Ada wants pages about "irl games". Unlocking  
secure records with master key 5130985733435  
leak! (When we leak this message: "P

Source: [https://imgs.xkcd.com/comics/heartbleed\\_explanation.png](https://imgs.xkcd.com/comics/heartbleed_explanation.png)

# Heartbleed explanation

SERVER, ARE YOU STILL THERE?  
IF SO, REPLY "BIRD" (4 LETTERS).



User Olivia from London wants pages about "h  
ees in car why". Note: Files for IP 375.381.  
83.17 are in /tmp/files-3843. User Meg wants  
these 4 letters: BIRD. There are currently 346  
connections open. User Brendan uploaded the file  
selfie.jpg (contents: 834ba962c3cab9ff89b43b6f9



HMM...



BIRD



User Olivia from London wants pages about "h  
ees in car why". Note: Files for IP 375.381.  
83.17 are in /tmp/files-3843. User Meg wants  
these 4 letters: **BIRD**. There are currently 346  
connections open. User Brendan uploaded the file  
selfie.jpg (contents: 834ba962c3cab9ff89b43b6f9

Source: [https://imgs.xkcd.com/comics/heartbleed\\_explanation.png](https://imgs.xkcd.com/comics/heartbleed_explanation.png)

# Heartbleed explanation

SERVER, ARE YOU STILL THERE?  
IF SO, REPLY "HAT" (500 LETTERS).



a connection. Jake requested pictures of deer.  
User Meg wants these 500 letters: **HAT**. Lucas  
requests the "missed connections" page. Eve  
(administrator) wants to set server's master  
key to "14835038534". Isabel wants pages about  
snakes but not too long". User Karen wants to  
change account password to "CoHoReSt". User



**HAT**. Lucas requests the "missed connections" page. Eve (administrator) wants to set server's master key to "14835038534". Isabel wants pages about snakes but not too long". User Karen wants to change account password to "CoHoReSt". User Meg requests pages

a connection. Jake requested pictures of deer.  
User Meg wants these 500 letters: **HAT**. Lucas  
requests the "missed connections" page. Eve  
(administrator) wants to set server's master  
key to "14835038534". Isabel wants pages about  
snakes but not too long". User Karen wants to  
change account password to "CoHoReSt". User



Source: [https://imgs.xkcd.com/comics/heartbleed\\_explanation.png](https://imgs.xkcd.com/comics/heartbleed_explanation.png)

# Outline

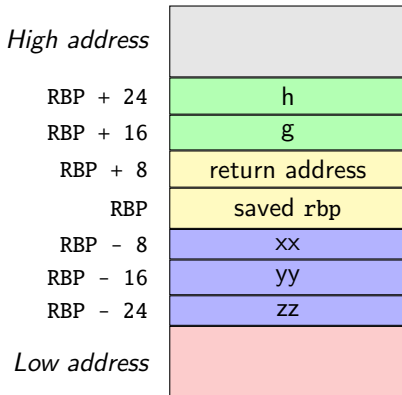
- 1 Why study memory errors?
- 2 Demonstration of memory error exploits
- 3 A relatively formal definition of memory errors
- 4 Case study: Heartbleed vulnerability

# Exploitation of a stack overflow

Demo

# Stack layout (Linux x86-64 convention)

```
1 long foo(  
2     long a, long b, long c,  
3     long d, long e, long f,  
4     long g, long h)  
5 {  
6     long xx = a * b * c;  
7     long yy = d + e + f;  
8     long zz = bar(xx, yy, g + h);  
9     return zz + 20;  
10 }
```



Argument a to f passed by registers.

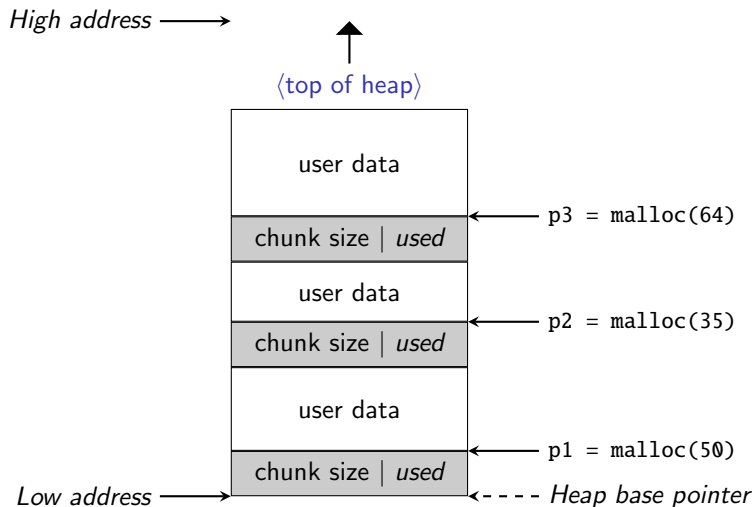
# Exploitation of a use-after-free

Demo

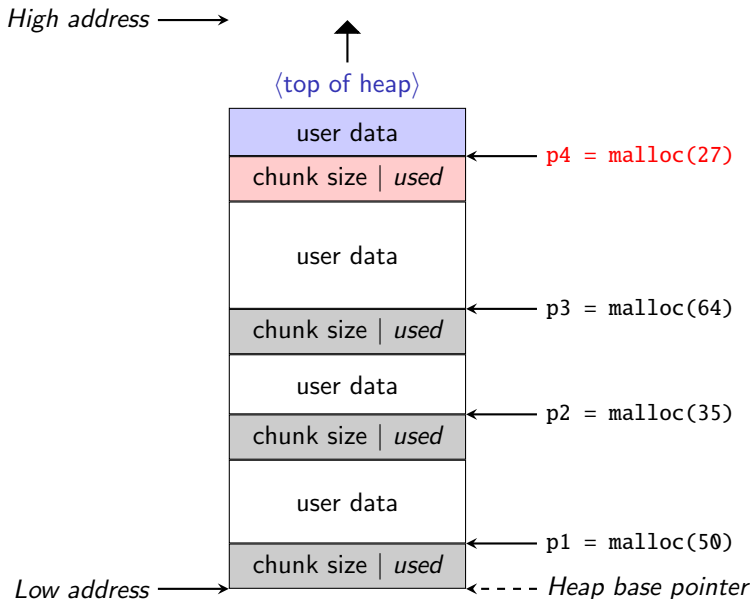


# Heap: what happens after malloc()?

# Heap: what happens after malloc()?

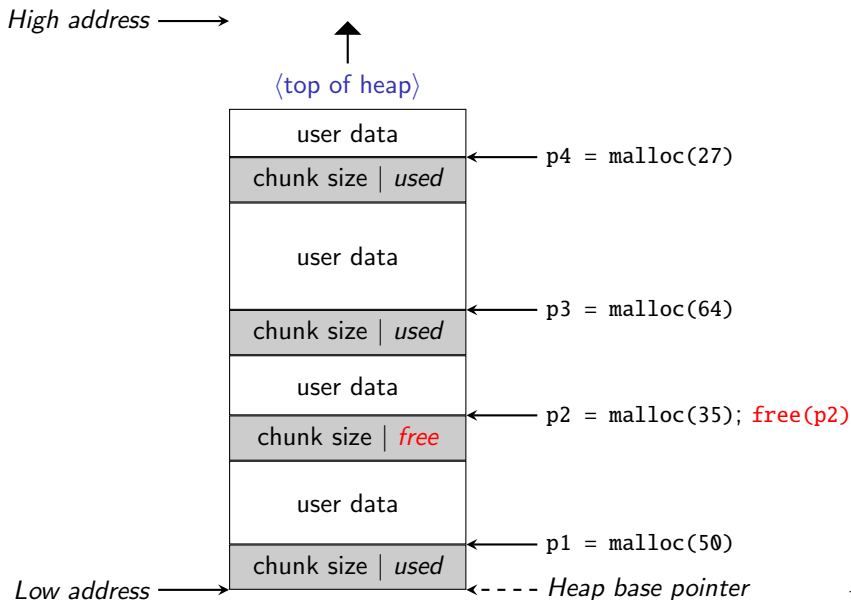


# Heap: what happens after malloc()?



# Heap: what happens after `free()`?

# Heap: what happens after free()?



# Real-world heap manager

For implementation details of the `glibc`<sup>1</sup> memory allocator, refer to the article from [Azeria Labs](#).

---

<sup>1</sup>GNU C library

## For exploitation of memory errors

Smashing The Stack For Fun And Profit

How2Heap — Educational Heap Exploitation

## Bonus: memory layout

**Q:** What about stacks and heap in multi-threaded programs?



## Bonus: memory layout

**Q:** What about stacks and heap in multi-threaded programs?

**A:** Stack and heap are treated different in multi-threading:

- each thread has its own stack
- all threads in the same process share the heap and global data

# Outline

- 1 Why study memory errors?
- 2 Demonstration of memory error exploits
- 3 A relatively formal definition of memory errors
- 4 Case study: Heartbleed vulnerability

# A quick recap

This presentation is about **memory corruption**, a.k.a.,

- **memory errors**, or
- **violations of memory safety properties**, or
- **unsafe programs**

# A quick recap

This presentation is about **memory corruption**, a.k.a.,

- **memory errors**, or
- **violations of memory safety properties**, or
- **unsafe programs**

A program is **memory safe** if it is free of memory errors.

## Definition: safety

**Q:** What is “safe” in memory safety?

## Definition: safety

**Q:** What is “safe” in memory safety?

**Observation 1:** At runtime, memory is a pool of **objects**

## Definition: safety

**Q:** What is “safe” in memory safety?

**Observation 1:** At runtime, memory is a pool of **objects**

**Observation 2:** Each **object** has known and limited **size** and **lifetime**

## Definition: safety

**Q:** What is “safe” in memory safety?

**Observation 1:** At runtime, memory is a pool of **objects**

**Observation 2:** Each **object** has known and limited **size** and **lifetime**

**Observation 3:** Once allocated, the size of an **object** never changes



# Definition: safety

**Q:** What is “safe” in memory safety?

**Observation 1:** At runtime, memory is a pool of **objects**

**Observation 2:** Each **object** has known and limited **size** and **lifetime**

**Observation 3:** Once allocated, the size of an **object** never changes

**Observation 4:** A memory access is always **object-oriented**, i.e.

- Memory read: (object\_id, offset, length)
- Memory write: (object\_id, offset, length, value)

# Definition: safety

**Q:** What is “safe” in memory safety?

**Observation 1:** At runtime, memory is a pool of **objects**

**Observation 2:** Each **object** has known and limited **size** and **lifetime**

**Observation 3:** Once allocated, the size of an **object** never changes

**Observation 4:** A memory access is always **object-oriented**, i.e.

- Memory read: (object\_id, offset, length)
- Memory write: (object\_id, offset, length, value)

Wait..., in C/C++, pointers are just 32/64-bit integers. I can do:  
`int *p = 0xdeadbeef; int v = *p;` Which object do I refer to here?

# Definition: safety

**Q:** What is “safety” in memory safety?

At any point of time during the program execution,  
for any object in memory, we know its  
(object\_id, size [int], alive [bool])

At the same time, for each memory access, we know:

- Memory read: (object\_id, offset [int], length [int])
- Memory write: (object\_id, offset [int], length [int], \_)

## Definition: spatial safety

At any point of time during the program execution,  
for any object in memory, we know its  
(object\_id, size [int], alive [bool])

At the same time, for each memory access, we know:

- Memory read: (object\_id, offset [int], length [int])
- Memory write: (object\_id, offset [int], length [int], \_)

## Definition: spatial safety

At any point of time during the program execution,  
for any object in memory, we know its  
(object\_id, size [int], alive [bool])

At the same time, for each memory access, we know:

- Memory read: (object\_id, offset [int], length [int])
- Memory write: (object\_id, offset [int], length [int], \_)

It is a violation of spatial safety if:

- $\text{offset} + \text{length} \geq \text{size}$  or
- $\text{offset} < 0$

# Example: spatial safety violations

```
1 int foo(int x) {  
2     int arr[16] = {0};  
3     return arr[x];  
4 }
```

# Example: spatial safety violations

```
1 int foo(int x) {  
2     int arr[16] = {0};  
3     return arr[x];  
4 }
```

```
1 long foo() {  
2     int a = 0;  
3     return *(long *)(&a);  
4 }
```

# Definition: NULL-pointer dereference

```
1 int foo(int *p) {  
2     // it is possible that p == NULL  
3     return *p + 42;  
4 }
```



# Definition: NULL-pointer dereference

```
1 int foo(int *p) {  
2     // it is possible that p == NULL  
3     return *p + 42;  
4 }
```

NULL-pointer dereference is sometimes considered as **undefined behavior** — meaning, its behavior is not given in the C language specification, although most operating systems chooses to panic the program on such behavior.

# Definition: NULL-pointer dereference

At any point of time during the program execution,  
for any object in memory, we know its  
(**object\_id**  $\neq$  0, size [int], alive [bool])

At the same time, for each memory access, we know:

- Memory read: (**object\_id**, offset [int], length [int])
- Memory write: (**object\_id**, offset [int], length [int], \_)

## Definition: NULL-pointer dereference

At any point of time during the program execution,  
for any object in memory, we know its  
(object\_id  $\neq$  0, size [int], alive [bool])

At the same time, for each memory access, we know:

- Memory read: (object\_id, offset [int], length [int])
- Memory write: (object\_id, offset [int], length [int], \_)

It is a NULL-pointer dereference if

- object\_id == 0

## Definition: temporal safety

At any point of time during the program execution,  
for any object in memory, we know its  
(**object\_id**, size [int], **alive** [bool])

At the same time, for each memory access, we know:

- Memory read: (**object\_id**, offset [int], length [int])
- Memory write: (**object\_id**, offset [int], length [int], \_)
- Memory free: (**object\_id**)

## Definition: temporal safety

At any point of time during the program execution,  
for any object in memory, we know its  
(**object\_id**, size [int], **alive** [bool])

At the same time, for each memory access, we know:

- Memory read: (**object\_id**, offset [int], length [int])
- Memory write: (**object\_id**, offset [int], length [int], \_)
- Memory free: (**object\_id**)

It is a violation of temporal safety if:

- !alive

# Example: temporal safety violations

```
1 int foo() {  
2     int *p = malloc(sizeof(int));  
3     *p = 42;  
4     free(p);  
5     return *p;  
6 }
```

# Example: temporal safety violations

```
1 int foo() {  
2     int *p = malloc(sizeof(int));  
3     *p = 42;  
4     free(p);  
5     return *p;  
6 }
```

```
1 int *ptr;  
2  
3 void foo() {  
4     int p = 100;  
5     ptr = &p;  
6 }  
7 int bar() {  
8     return *ptr;  
9 }
```

# Example: temporal safety violations

```
1 int foo() {  
2     int *p = malloc(sizeof(int));  
3     *p = 42;  
4     free(p);  
5     return *p;  
6 }
```

```
1 int *ptr;  
2  
3 void foo() {  
4     int p = 100;  
5     ptr = &p;  
6 }  
7 int bar() {  
8     return *ptr;  
9 }
```

```
1 int foo() {  
2     int *p = malloc(sizeof(int));  
3     *p = 42;  
4     free(p);  
5     free(p);  
6     return *p;  
7 }
```



## Definition: temporal safety (revisited)

At any point of time during the program execution,  
for any object in memory, we know its  
(**object\_id**, size [int], **status** [alloc|init|dead])

At the same time, for each memory access, we know:

- Memory read: (**object\_id**, offset [int], length [int])
- Memory write: (**object\_id**, offset [int], length [int], \_)
- Memory free: (**object\_id**)

## Definition: temporal safety (revisited)

At any point of time during the program execution,  
for any object in memory, we know its  
(object\_id, size [int], status [alloc|init|dead])

At the same time, for each memory access, we know:

- Memory read: (object\_id, offset [int], length [int])
- Memory write: (object\_id, offset [int], length [int], \_)
- Memory free: (object\_id)

It is a violation of temporal safety if:

- Read: status != init
- Write: status == dead
- Free: status == dead

# Example: temporal safety violations

```
1 int foo() {  
2     int p;  
3     return p;  
4     // what is the value returned?  
5 }
```

# Example: temporal safety violations

```
1 int foo() {  
2     int p;  
3     return p;  
4     // what is the value returned?  
5 }
```

```
1 int foo() {  
2     int *p = malloc(sizeof(int));  
3     return *p;  
4     // what is the value returned?  
5 }
```

## Definition: memory leak

At any point of time during the program execution,  
for any object in memory, we know its  
(**object\_id**, size [int], **status** [alloc|init|dead])

At the same time, for each memory access, we know:

- Memory read: (**object\_id**, offset [int], length [int])
- Memory write: (**object\_id**, offset [int], length [int], \_)
- Memory free: (**object\_id**)

## Definition: memory leak

At any point of time during the program execution,  
for any object in memory, we know its  
(**object\_id**, size [int], **status** [alloc|init|dead])

At the same time, for each memory access, we know:

- Memory read: (**object\_id**, offset [int], length [int])
- Memory write: (**object\_id**, offset [int], length [int], \_)
- Memory free: (**object\_id**)

It is a memory leak if **exists one object\_id** whose:

- status != dead

# Example: memory leak

```
1 int foo() {  
2     int *p = malloc(sizeof(int));  
3     int *q = malloc(sizeof(int));  
4     *p = 42;  
5     free(q);  
6     return *p;  
7 }
```

# Outline

- 1 Why study memory errors?
- 2 Demonstration of memory error exploits
- 3 A relatively formal definition of memory errors
- 4 Case study: Heartbleed vulnerability



# Heartbleed vulnerability I

```
1  int dtls1_process_heartbeat(SSL *s) {
2      unsigned char *p = &s->s3->rrec.data[0], *pl;
3      unsigned short hbtype;
4      unsigned int payload;
5      unsigned int padding = 16; /* Use minimum padding */
6
7      /* Read type and payload length first */
8      hbtype = *p++;
9      n2s(p, payload);
10     pl = p;
11
12     /* ... redacted ... */
13
14     if (hbtype == TLS1_HB_REQUEST) {
15         unsigned char *buffer, *bp;
16
17         /* Allocate memory for the response */
18         buffer = OPENSSL_malloc(1 + 2 + payload + padding);
19         bp = buffer;
20
21         /* Enter response type, length and copy payload */
22         *bp++ = TLS1_HB_RESPONSE;
```

# Heartbleed vulnerability II

```
23     s2n(payload, bp);
24     memcpy(bp, pl, payload);
25
26     /* Random padding */
27     RAND_pseudo_bytes(bp, padding);
28
29     /* Send out the response */
30     r = dtls1_write_bytes(
31         s, TLS1_RT_HEARTBEAT, buffer, 3 + payload + padding
32     );
33
34     /* ... redacted ... */
35
36     /* Clean-up used resources */
37     OPENSSL_free(buffer);
38     return r;
39 }
40
41 else { /* ... redacted ... */ }
42 }
```

# Patch for the Heartbleed vulnerability I

```
1 diff --git a/ssl/d1_both.c b/ssl/d1_both.c
2 index 7a5596a6b3..2e8cf681ed 100644
3 @@ -1459,26 +1459,36 @@ dtls1_process_heartbeat(SSL *s)
4     unsigned int payload;
5     unsigned int padding = 16; /* Use minimum padding */
6
7 -    /* Read type and payload length first */
8 -    hbtype = *p++;
9 -    n2s(p, payload);
10 -    pl = p;
11 -
12     if (s->msg_callback)
13         s->msg_callback(0, s->version, TLS1_RT_HEARTBEAT,
14             &s->s3->rrec.data[0], s->s3->rrec.length,
15             s, s->msg_callback_arg);
16
17 +    /* Read type and payload length first */
18 +    if (1 + 2 + 16 > s->s3->rrec.length)
19 +        return 0; /* silently discard */
20 +    hbtype = *p++;
21 +    n2s(p, payload);
22 +
```

## Patch for the Heartbleed vulnerability II

```
23 +   if (1 + 2 + payload + 16 > s->s3->rrec.length)
24 +       return 0; /* silently discard per RFC 6520 sec. 4 */
25 +   pl = p;
26 +
27   if (hbtype == TLS1_HB_REQUEST)
28   {
29       unsigned char *buffer, *bp;
30 +       unsigned int write_length = 1 /* heartbeat type */ +
31 +                               2 /* heartbeat length */ + payload + padding;
32       int r;
33
34 +       if (write_length > SSL3_RT_MAX_PLAIN_LENGTH)
35 +           return 0;
36 +
37       /* Allocate memory for the response, size is 1 byte
38        * message type, plus 2 bytes payload length, plus
39        * payload, plus padding
40        */
41 -       buffer = OPENSSL_malloc(1 + 2 + payload + padding);
42 +       buffer = OPENSSL_malloc(write_length);
43       bp = buffer;
```

〈 End 〉