

CS 453/698: Software and Systems Security

Module: An In-depth Study of Memory Errors

Lecture: Exploit mitigation

Meng Xu (*University of Waterloo*)

Winter 2025

Outline

- 1 Introduction: what is mitigation?
- 2 Principle of least privileges (PoLP)
- 3 Reference monitoring

Software security landscape

Generally speaking, almost all work in the software security area can be categorized into four bins:

- **Vulnerability:** *Identify a bug in the program that may cause some damage*
 - $f(\text{Code}) \rightarrow \text{Bug}$
- **Exploitation:** *Given a set of bugs, exploit them to achieve a desired goal*
 - $f(\text{Code}, \{\dots\text{Bug}\dots\}, \text{Goal}) \rightarrow \text{Action}$
- **Mitigation:** *Given a set of bugs and an associated set of exploits, prevent them*
 - $f(\text{Code}, \{\dots\text{Bug}\dots\}, \{\dots\text{Action}\dots\}) \rightarrow \text{Blockage}$
- **Detection:** *Given a program, check the existence of a specific type of bug*
 - $f(\text{Code}, \text{Bug}, [\text{Action}]) \rightarrow \text{Signal}$
- **Prevention:** *It is impossible to create a program that has a specific type of bug*

Themes of mitigation

- Principle of least privileges (PoLP)
- Reference monitoring / program shepherding
- Moving-target defense

Themes of mitigation

- Principle of least privileges (PoLP)
 - reduce permissions unless absolutely needed
- Reference monitoring / program shepherding
- Moving-target defense

Themes of mitigation

- Principle of least privileges (PoLP)
 - reduce permissions unless absolutely needed
- Reference monitoring / program herding
 - keep an eye on the program while it is executing
- Moving-target defense

Themes of mitigation

- Principle of least privileges (PoLP)
 - reduce permissions unless absolutely needed
- Reference monitoring / program shepherding
 - keep an eye on the program while it is executing
- Moving-target defense
 - non-determinism is useful in software security when
 - * it has no impact on the intended finite state machine BUT
 - * limits attackers' abilities to program the weird machine.

Outline

- 1 Introduction: what is mitigation?
- 2 Principle of least privileges (PoLP)
- 3 Reference monitoring

DEP a.k.a., $W \oplus X$

DEP a.k.a., $W \oplus X$

DEP – Data Execution Prevention

$W \oplus X$ – Write exclusive-or eXecute

You can either **write data** **OR** **execute code** in a memory region,
but **never both**.

DEP a.k.a., $W \oplus X$

DEP – Data Execution Prevention

$W \oplus X$ – Write exclusive-or eXecute

You can either **write data** **OR** **execute code** in a memory region,
but **never both**.

Implementation: `gcc -z execstack`.

Motivation for type-based heap allocation

Motivation for type-based heap allocation

A more realistic use-after-free (UAF) exploit:

```
1 struct N {  
2     long user;  
3     int (*fn)(void);  
4 };
```

```
1 struct O {  
2     int (*oper)(void);  
3     long id;  
4 };
```

```
1 void foo(long user) {  
2     struct N *p =  
3         malloc(sizeof(struct N));  
4  
5     p->fn = __safe_function_1;  
6     p->user = user;  
7  
8     /* ... */  
9     /* later in the code */  
10    /* ... */  
11    p->fn();  
12 }
```

```
1 void bar(long id) {  
2     struct O *x =  
3         malloc(sizeof(struct O));  
4  
5     x->oper = __safe_function_2;  
6     x->id = id;  
7     struct O *q = x;  
8     free(x);           // q is dangling  
9  
10    /* later in the code */  
11    q->oper();  
12 }
```

Sample UAF-exploit (continued)

```
1 {
2     /* from bar(..) */
3     struct 0 *x =
4         malloc(sizeof(struct 0));
5
6     x->oper = __safe_function_2;
7     x->id = id;
8     struct 0 *q = x;
9     free(x);           // q is dangling
10
11     /* from foo(..) */
12     struct N *p =
13         malloc(sizeof(struct N));
14
15     p->fn = __safe_function_1;
16     p->user = user;
17
18     /* from bar(..) */
19     q->oper();
20 }
```

Type-based heap allocation

If a memory address refers to a heap object of type `T`, it will **always** refer to objects of type `T`, no matter what (e.g., freed and re-allocated).

NOTE: this does not imply that this memory address will be assigned to a `T *` pointer. It can be assigned to a `void *`, an `int *`, or anything.

Outline

- 1 Introduction: what is mitigation?
- 2 Principle of least privileges (PoLP)
- 3 Reference monitoring

CFI: introduction

Control-Flow Integrity (CFI) is a classic example of **runtime reference monitor** in software security.

CFI: introduction

Control-Flow Integrity (CFI) is a classic example of **runtime reference monitor** in software security.

CFI is also sometimes referred to as **program shepherding**

monitoring control flow transfers during program execution to enforce a security policy — from [a paper in USENIX Security'02](#).

Basic use cases of CFI

```
1 void f1();
2 void f2();
3 void f3();
4 void f4(int, int);
5
6 void foo(int usr) {
7     void (*func)();
8
9     if (usr == MAGIC)
10         func = f1;
11     else
12         func = f2;
13
14     // forward edge CFI check
15     CHECK_CFI_FORWARD(func);
16     func();
17
18     // backward edge CFI check
19     CHECK_CFI_BACKWARD();
20 }
```

Basic use cases of CFI

Option 1: allow all functions

- f1, f2, f3, f4, foo, printf, system, ...

```
1 void f1();
2 void f2();
3 void f3();
4 void f4(int, int);
5
6 void foo(int usr) {
7     void (*func)();
8
9     if (usr == MAGIC)
10         func = f1;
11     else
12         func = f2;
13
14     // forward edge CFI check
15     CHECK_CFI_FORWARD(func);
16     func();
17
18     // backward edge CFI check
19     CHECK_CFI_BACKWARD();
20 }
```

Basic use cases of CFI

```
1 void f1();
2 void f2();
3 void f3();
4 void f4(int, int);
5
6 void foo(int usr) {
7     void (*func)();
8
9     if (usr == MAGIC)
10         func = f1;
11     else
12         func = f2;
13
14     // forward edge CFI check
15     CHECK_CFI_FORWARD(func);
16     func();
17
18     // backward edge CFI check
19     CHECK_CFI_BACKWARD();
20 }
```

Option 1: allow all functions

- f1, f2, f3, f4, foo, printf, system, ...

Option 2: allowed only functions defined in the current module

- f1, f2, f3, f4, foo

Basic use cases of CFI

```
1 void f1();
2 void f2();
3 void f3();
4 void f4(int, int);
5
6 void foo(int usr) {
7     void (*func)();
8
9     if (usr == MAGIC)
10         func = f1;
11     else
12         func = f2;
13
14     // forward edge CFI check
15     CHECK_CFI_FORWARD(func);
16     func();
17
18     // backward edge CFI check
19     CHECK_CFI_BACKWARD();
20 }
```

Option 1: allow all functions

- f1, f2, f3, f4, foo, printf, system, ...

Option 2: allowed only functions defined in the current module

- f1, f2, f3, f4, foo

Option 3: allow functions with type signature `void (*)()`

- f1, f2, f3

Basic use cases of CFI

```
1 void f1();
2 void f2();
3 void f3();
4 void f4(int, int);
5
6 void foo(int usr) {
7     void (*func)();
8
9     if (usr == MAGIC)
10         func = f1;
11     else
12         func = f2;
13
14     // forward edge CFI check
15     CHECK_CFI_FORWARD(func);
16     func();
17
18     // backward edge CFI check
19     CHECK_CFI_BACKWARD();
20 }
```

Option 1: allow all functions

- f1, f2, f3, f4, foo, printf, system, ...

Option 2: allowed only functions defined in the current module

- f1, f2, f3, f4, foo

Option 3: allow functions with type signature `void (*)()`

- f1, f2, f3

Option 4: allow functions whose address are taken (e.g., assigned)

- f1, f2

Example: Microsoft Return-flow Guard (RFG)

RFG was our compatible, ABI compliant, performant software shadow stack

Compile Time

NOP's added to the prolog & epilog of all functions

Metadata added to the image to locate the prolog and epilog NOP bytes

Runtime

Process Start

- 1TB shadow stack region created
- Region cannot be queried
- A/V's in region are fatal
- FS segment points to the shadow stack of the current thread

Image Load

- If process enables RFG: patch NOP's with RFG prolog/epilog

Function Calls

- Prolog: Push return address to shadow stack
- Epilog: Fast fail if return address on stack and shadow stack are mismatched

Parent Function	Child Function
[...] //Prior code	
call ChildFunction	
	mov rax, [rsp]
	mov fs:[rsp], rax
	[...] //Child code
	mov rcx, fs:[rsp]
	cmp rcx, [rsp]
	jne _fast_fail
	ret
0xABCD: [...] //Remainder of parent function	

If attacker changes the return address at these points RFG is defeated

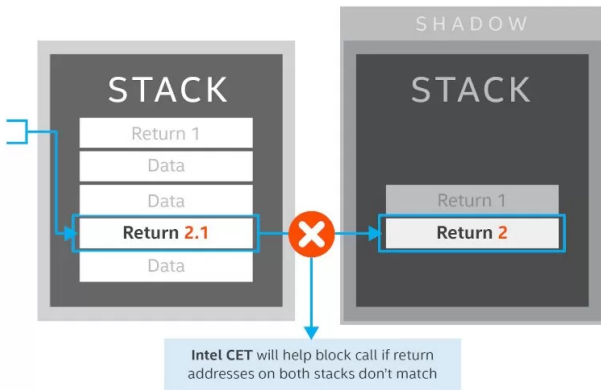
RFG relies on a secret: the shadow stack's virtual address

Illustration taken from [Microsoft Talk: The Evolution of CFI Attacks and Defenses](#)

Back-edge protection: shadow stack

SHADOW STACK (SS)

SS delivers return address protection to defend against return-oriented programming (ROP) attack methods.



CET: shadow stack

- For every regular stack CET adds a shadow stack region, which is indexed via a new register %ssp.
- Regular memory stores (executed from any ring) are not allowed in shadow stack region

When enabled,

- Each time a call instruction gets executed, in addition to the return address being pushed onto the regular stack, a copy of it is also pushed (automatically) onto the shadow stack.
- Each time a ret instruction gets executed, the return addresses pointed by %rsp and %ssp are (automatically) popped from the two stacks, and their values are compared together.

CET: Indirect Branch Tracking (IBT)

CET introduces a new (4-byte) instruction, i.e., `endbr`, which becomes the **only** allowed target of indirect `call/jmp` instructions.

In other words, forward-edge transfers via (indirect) `call` or `jmp` instructions are pinned to code locations that are “marked” with an `endbr`; else, an exception (`#CP`) is raised.

IBT example

```
1 void main() {  
2     int (*f) {};  
3     f = foo;  
4     f();  
5 }  
6  
7 int foo() {  
8     return 0;  
9 }
```

```
1 <main>:  
2 movq    $0x4004fb, -8(%rbp)  
3 mov     -8(%rbp), %rdx  
4 call    *%rdx  
5 :  
6 retq  
7  
8 <foo>:  
9 endbr64  
10 :  
11 mov     rax, 0  
12 :  
13 retq
```

IBT example

```
1 void main() {  
2     int (*f) {};  
3     int (*g) {};  
4     f = foo;  
5     g = bar;  
6     f();  
7     g();  
8 }  
9  
10 int foo() {  
11     return 0;  
12 }  
13  
14 int bar() {  
15     return 1;  
16 }
```

```
1 <main>:  
2 movq    $0x4004fb, -16(%rbp)  
3 mov     -16(%rbp), %rdx  
4 call    *%rdx  
5 mov     -8(%rbp), %rdx  
6 call    *%rdx  
7 :  
8 retq  
9  
10 <foo>:  
11 endbr64  
12 :  
13 mov     rax, 0  
14 :  
15 retq  
16  
17 <bar>:  
18 endbr64  
19 :  
20 mov     rax, 1  
21 :  
22 retq
```

Security boundaries of CFI-protected programs

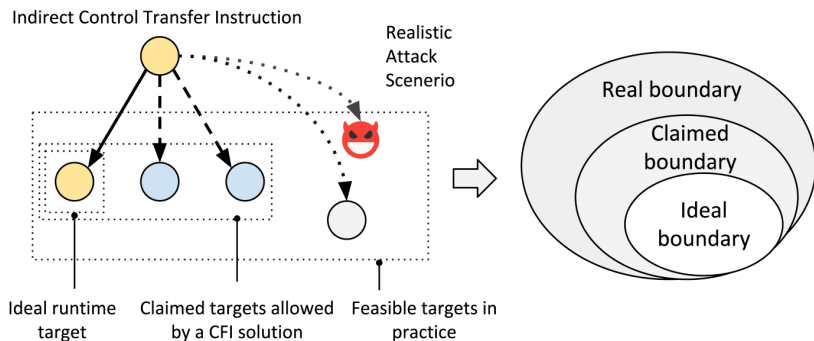


Figure from [a paper published in ACM CCS'20](#)

Pointer integrity

Goal: ensures **pointers** in memory remain **unchanged**.

Pointer integrity

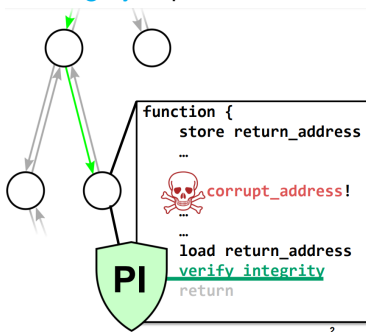
Goal: ensures **pointers** in memory remain **unchanged**.

- i.e., the **value of the pointer** remains unchanged, not the memory content referred to by this pointer.

Pointer integrity

Goal: ensures **pointers** in memory remain **unchanged**.

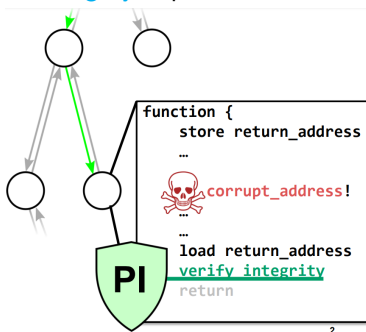
- i.e., the **value of the pointer** remains unchanged, not the memory content referred to by this pointer.
- Perfect **code pointer integrity** implies control-flow integrity (CFI).



Pointer integrity

Goal: ensures **pointers** in memory remain **unchanged**.

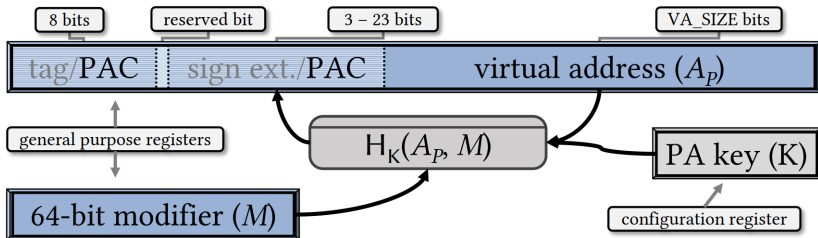
- i.e., the **value of the pointer** remains unchanged, not the memory content referred to by this pointer.
- Perfect **code pointer integrity** implies control-flow integrity (CFI).



- Data pointer integrity is also important (e.g., against data-only attacks and data-oriented programming) and can be (partially) achieved via Pointer Authentication.

Overview of Arm Pointer Authentication (PA)

Available since Armv8.3-A instruction set architecture (ISA) when the processor executes in 64-bit Arm state (AArch64)



PA consists of a set of instructions for creating and authenticating **pointer authentication codes (PACs)**.

PAC details

- Each PAC is derived from
 - A pointer value
 - A 64-bit context value (modifier)
 - A 128-bit secret key

PAC details

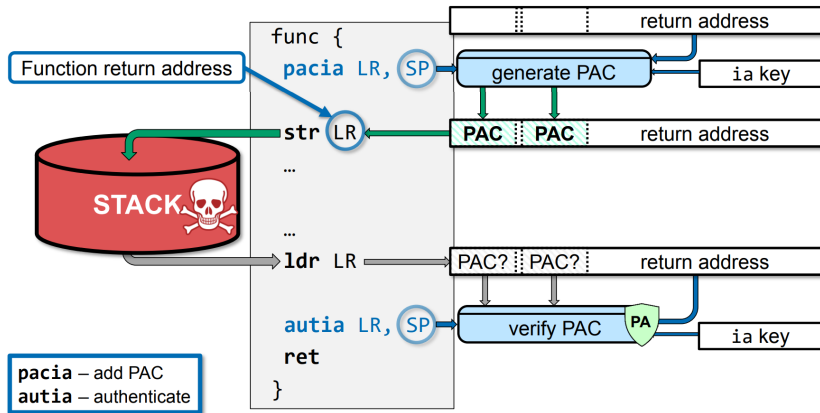
- Each PAC is derived from
 - A pointer value
 - * an N-bit memory address
 - A 64-bit context value (modifier)
 - * doesn't need to secret, as long as it provides enough entropy
 - A 128-bit secret key
 - * held in system registers, set by the kernel per each process,
 - * can be used, but cannot be read/written by userspace

PAC details

- Each PAC is derived from
 - A pointer value
 - * an N-bit memory address
 - A 64-bit context value (modifier)
 - * doesn't need to secret, as long as it provides enough entropy
 - A 128-bit secret key
 - * held in system registers, set by the kernel per each process,
 - * can be used, but cannot be read/written by userspace
- PAC essentially a key-ed message authentication code (MAC) where the MAC algorithm can be implementation defined
 - by default, it is **QARMA**
- Instructions hide the algorithm details (sign + authenticate)

Example: PA-based return address signing

Deployed as `-msign-return-address` in GCC and LLVM/Clang



〈 End 〉