

# CS 453/698: Software and Systems Security

## **Module: Other Common Vulnerability Types**

Lecture: Common mistakes

Meng Xu (*University of Waterloo*)

Winter 2025

# Outline

- 1 Introduction: why study these bug types?
- 2 Undefined / counterintuitive behaviors
- 3 Insufficient sanitization on untrusted input
- 4 Invocation of / by untrusted logic
- 5 Hostile execution environment
- 6 Conclusion

## “Nice” properties of memory errors

- They have **universally** accepted definitions
  - Once you find a memory error, you do not need to diligently argue that this is a bug and not a feature
- They often lead to a set of known consequences that are **generally** considered severe (e.g., data leak or denial-of-service)
  - Once you find a memory error, you do not need to construct a working exploit to justify it
- Finding them typically **do not require** program-specific domain knowledge
  - If you have a technique that can find memory errors in one codebase, you can scale it up to millions of codebases

## “Nice” properties of memory errors

- They have **universally** accepted definitions
  - Once you find a memory error, you do not need to diligently argue that this is a bug and not a feature
- They often lead to a set of known consequences that are **generally** considered severe (e.g., data leak or denial-of-service)
  - Once you find a memory error, you do not need to construct a working exploit to justify it
- Finding them typically **do not require** program-specific domain knowledge
  - If you have a technique that can find memory errors in one codebase, you can scale it up to millions of codebases

In fact, very few types of vulnerabilities meet these requirements.

## “Nice” properties of memory errors

- They have **universally** accepted definitions
  - Once you find a memory error, you do not need to diligently argue that this is a bug and not a feature
- They often lead to a set of known consequences that are **generally** considered severe (e.g., data leak or denial-of-service)
  - Once you find a memory error, you do not need to construct a working exploit to justify it
- Finding them typically **do not require** program-specific domain knowledge
  - If you have a technique that can find memory errors in one codebase, you can scale it up to millions of codebases

In fact, very few types of vulnerabilities meet these requirements.

⇒ Most of the bug types covered today **do not** meet all requirements, but they are representative examples to show easy it is to make a mistake in programming.

# Outline

- 1 Introduction: why study these bug types?
- 2 Undefined / counterintuitive behaviors**
- 3 Insufficient sanitization on untrusted input
- 4 Invocation of / by untrusted logic
- 5 Hostile execution environment
- 6 Conclusion

# Unsafe integer operations

Mathematical integers are **unbounded**

## WHILE

Machine integers are **bounded** by a fixed number of bits.

# Unsafe integer operations

```
1 mapping (address => uint256) public balanceOf;
2
3 // INSECURE
4 function transfer(address _to, uint256 _value) {
5     /* Check if sender has balance */
6     require(balanceOf[msg.sender] >= _value);
7
8     /* Add and subtract new balances */
9     balanceOf[msg.sender] -= _value;
10    balanceOf[_to] += _value;
11 }
```

**Q:** What is the bug here?



# Unsafe integer operations

```
1 mapping (address => uint256) public balanceOf;
2
3 // INSECURE
4 function transfer(address _to, uint256 _value) {
5     /* Check if sender has balance */
6     require(balanceOf[msg.sender] >= _value);
7
8     /* Add and subtract new balances */
9     balanceOf[msg.sender] -= _value;
10    balanceOf[_to] += _value;
11 }

1 // SECURE
2 function transfer(address _to, uint256 _value) {
3     /* Check if sender has balance and for overflows */
4     require(balanceOf[msg.sender] >= _value &&
5             balanceOf[_to] + _value >= balanceOf[_to]);
6
7     /* Add and subtract new balances */
8     balanceOf[msg.sender] -= _value;
9     balanceOf[_to] += _value;
10 }
```

# Common cases for integer overflows and underflows

- signed  $\leftrightarrow$  unsigned
- size-decreasing cast (a.k.a., truncate)
- +, -, \* for both signed and unsigned integers
- / for signed integers
- ++ and -- for both signed and unsigned integers
- +=, -=, \*= for both signed and unsigned integers
- /= for signed integers
- Negation - for signed and unsigned integers
- << for both signed and unsigned integers

# Unsafe floating-point operations

Mathematical real numbers are arbitrary precision

## WHILE

Machine floating-point numbers are bounded by a limited precision.

# The perils of floating point (in Python)

```
>>> .1 + .1 + .1 == .3
```

**Q:** True or False?

# The perils of floating point (in Python)

```
>>> .1 + .1 + .1 == .3
```

**Q:** True or False?

```
>>> round(.1, 1) + round(.1, 1) + round(.1, 1) == round(.3, 1)
```

**Q:** True or False?

# The perils of floating point (in Python)

```
>>> .1 + .1 + .1 == .3
```

**Q:** True or False?

```
>>> round(.1, 1) + round(.1, 1) + round(.1, 1) == round(.3, 1)
```

**Q:** True or False?

```
>>> round(.1 + .1 + .1, 10) == round(.3, 10)
```

**Q:** True or False?

# The perils of floating point (in Python)

```
>>> .1 + .1 + .1 == .3
```

**Q:** True or False?

```
>>> round(.1, 1) + round(.1, 1) + round(.1, 1) == round(.3, 1)
```

**Q:** True or False?

```
>>> round(.1 + .1 + .1, 10) == round(.3, 10)
```

**Q:** True or False?

Further reading: [The Perils of Floating Point](#)

# Pointer relational comparison

```
1 #include <stdio.h>
2
3 struct Record {
4     int a;
5     int b;
6 };
7
8 int main(void) {
9     struct Record r = { 0, 0 };
10    /* defined behavior */
11    if (&r.a < &r.b) {
12        printf("Hello\n");
13    } else {
14        printf("World\n");
15    }
16    return 0;
17 }
```

**Q:** Output?



# Pointer relational comparison

```
1 #include <stdio.h>
2
3 struct Record {
4     int a;
5     int b;
6 };
7
8 int main(void) {
9     struct Record r = { 0, 0 };
10    /* defined behavior */
11    if (&r.a < &r.b) {
12        printf("Hello\n");
13    } else {
14        printf("World\n");
15    }
16    return 0;
17 }
```

```
1 #include <stdio.h>
2
3 int main(void) {
4     int a = 0;
5     int b = 0;
6     /* undefined behavior */
7     if (&a < &b) {
8         printf("Hello\n");
9     } else {
10        printf("World\n");
11    }
12    return 0;
13 }
```

Q: Output?

Q: Output?

# Pointer relational comparison

In C and C++, the **relational comparison of pointers** to objects (i.e.,  $<$  or  $>$ ) is only strictly defined if

- the pointers point to **members of the same object**, or
- the pointers point to **elements of the same array**.

# Pointer relational comparison

In C and C++, the **relational comparison of pointers** to objects (i.e.,  $<$  or  $>$ ) is only strictly defined if

- the pointers point to **members of the same object**, or
- the pointers point to **elements of the same array**.

However, most compilers will emit a comparison operation based on the numerical value of the pointers.

# Pointer relational comparison

In C and C++, the **relational comparison of pointers** to objects (i.e.,  $<$  or  $>$ ) is only strictly defined if

- the pointers point to **members of the same object**, or
- the pointers point to **elements of the same array**.

However, most compilers will emit a comparison operation based on the numerical value of the pointers.  $\implies$  This is not strictly a bug, as **undefined behavior** means the compiler is free to choose whatever action that might make sense.

# Outline

- 1 Introduction: why study these bug types?
- 2 Undefined / counterintuitive behaviors
- 3 Insufficient sanitization on untrusted input**
- 4 Invocation of / by untrusted logic
- 5 Hostile execution environment
- 6 Conclusion

# Untrusted input

Handling untrusted input can be **dangerous!**

# SQL injection

```
1 public boolean login(String username, String password) {
2     String sql =
3         "SELECT * FROM Users WHERE " +
4         "username = '" + username + "' AND " +
5         "password = '" + password + "';";
6
7     ResultSet result = db.executeQuery(sql);
8     if (result.next()) {
9         /* login success */
10        return true;
11    } else {
12        /* login failure */
13        return false;
14    }
15 }
```

# Mitigating SQL injection with sanitization

```
1 public boolean login(String username, String password) {
2     PreparedStatement sql = db.prepareStatement(
3         "SELECT * FROM Users WHERE username = ? AND password = ?;")
4     sql.setString(1, username);
5     sql.setString(2, password);
6
7     ResultSet result = db.executeQuery(sql);
8     if (result.next()) {
9         /* login success */
10        return true;
11    } else {
12        /* login failure */
13        return false;
14    }
15 }
```



# SQL injection in the wild




Original source unknown, found on Twitter

# printf is powerful

A format string vulnerability is a bug where **untrusted user input** is passed as the format argument to `printf`, `scanf`, or another function in that family.

For details, see the [man page of printf](#).

# printf is powerful



# Format string vulnerability demo

```
1 #include <stdio.h>
2 #include <unistd.h>
3
4 int main() {
5     int secret = 0xdeadbeef;
6
7     char name[64] = {0};
8     read(0, name, 64);
9     printf("Hello ");
10    printf(name);
11    printf(", try to get the secret!\n");
12    return 0;
13 }
```

# Format string vulnerability demo

```
1 #include <stdio.h>
2 #include <unistd.h>
3
4 int main() {
5     int secret = 0xdeadbeef;
6
7     char name[64] = {0};
8     read(0, name, 64);
9     printf("Hello ");
10    printf(name);
11    printf(", try to get the secret!\n");
12    return 0;
13 }
```

To trigger the vulnerability, try something like `%7$11x`, although `%7` can be other values depending on the OS and C compiler version.

# Cross-site scripting (XSS)

Cross-site scripting (XSS) enables **attackers** to inject client-side scripts into **web pages** viewed by **other users**.

# Same-origin policy

This essentially states that if content from one site (such as <https://crisp.uwaterloo.ca>) is granted permission to access resources (e.g., cookies etc.) on a web browser, then content from **the same origin** will share these permissions.

# Same-origin policy

This essentially states that if content from one site (such as <https://crisp.uwaterloo.ca>) is granted permission to access resources (e.g., cookies etc.) on a web browser, then content from **the same origin** will share these permissions.

The same-origin property is defined as two URLs sharing the same

- URI scheme (e.g. ftp, http, or https)
- hostname (e.g., `crisp.uwaterloo.ca`) and
- port number (e.g., 80)



# Same-origin policy

This essentially states that if content from one site (such as <https://crisp.uwaterloo.ca>) is granted permission to access resources (e.g., cookies etc.) on a web browser, then content from **the same origin** will share these permissions.

The same-origin property is defined as two URLs sharing the same

- URI scheme (e.g. ftp, http, or https)
- hostname (e.g., [crisp.uwaterloo.ca](https://crisp.uwaterloo.ca)) and
- port number (e.g., 80)

For example, these webpages are from the same origin:

- <https://crisp.uwaterloo.ca/research/> and
- <https://crisp.uwaterloo.ca/courses/>

# XSS Demo I

```
1 from urllib.parse import unquote as url_unquote
2 from http.server import BaseHTTPRequestHandler, HTTPServer
3
4 HOST = "localhost"
5 PORT = 8080
6
7 PAGE = """<html>
8 <form action='/submit' method='POST'>
9 <input type='text' name='comment' />
10 </form>
11 </html>"""
12
13 class XSSDemoServer(BaseHTTPRequestHandler):
14     def do_GET(self):
15         self.send_response(200)
16         self.send_header("Content-type", "text/html")
17         self.end_headers()
18         self.wfile.write(bytes(PAGE, "utf-8"))
19
20     def do_POST(self):
21         size = int(self.headers.get('Content-Length'))
22         body = url_unquote(self.rfile.read(size).decode('utf-8'))
```

# XSS Demo II

```
23     self.send_response(200)
24     self.send_header("Content-type", "text/html")
25     self.end_headers()
26     self.wfile.write(bytes("<html>%s</html>" % body[8:], "utf-8"))
27
28
29 if __name__ == "__main__":
30     server = HTTPServer((HOST, PORT), XSSDemoServer)
31     print("Server started http://%s:%s" % (HOST, PORT))
32
33     try:
34         server.serve_forever()
35     except KeyboardInterrupt:
36         pass
37
38     server.server_close()
39     print("Server stopped.")
```

**Q:** Try `<script>alert("XSS")</script>`

# Outline

- 1 Introduction: why study these bug types?
- 2 Undefined / counterintuitive behaviors
- 3 Insufficient sanitization on untrusted input
- 4 Invocation of / by untrusted logic**
- 5 Hostile execution environment
- 6 Conclusion

# Calling into untrusted code is dangerous

# Calling into untrusted code is dangerous

## The DAO attack on Ethereum

# Calling into untrusted code is dangerous

## The DAO attack on Ethereum

*In 2016, an attacker exploited a **vulnerability** in The DAO's wallet smart contracts. In a couple of weeks (by Saturday, 18th June), the attacker managed to drain more than 3.6 million ether into an attacker-controlled account. The price of ether dropped from over \$20 to under \$13.*

# Calling into untrusted code is dangerous

## The DAO attack on Ethereum

*In 2016, an attacker exploited a **vulnerability** in The DAO's wallet smart contracts. In a couple of weeks (by Saturday, 18th June), the attacker managed to drain more than 3.6 million ether into an attacker-controlled account. The price of ether dropped from over \$20 to under \$13.*

The DAO attack was partially recovered by a **hard-fork** of the Ethereum blockchain that returns all stolen ethers into a special smart contract (which can be subsequently withdrawn). This resulted in two chains: Ethereum classic and Ethereum.



# Reentrancy attack (victim contract)

```
1 contract EtherStore {
2     uint256 public withdrawalLimit = 1 ether;
3     mapping(address => uint256) public lastWithdrawTime;
4     mapping(address => uint256) public balances;
5
6     function depositFunds() public payable {
7         balances[msg.sender] += msg.value;
8     }
9
10    function withdrawFunds (uint256 _weiToWithdraw) public {
11        require(balances[msg.sender] >= _weiToWithdraw);
12        require(_weiToWithdraw <= withdrawalLimit);
13        require(now >= lastWithdrawTime[msg.sender] + 1 weeks);
14        require(msg.sender.call.value(_weiToWithdraw)());
15
16        balances[msg.sender] -= _weiToWithdraw;
17        lastWithdrawTime[msg.sender] = now;
18    }
19 }
```

# Reentrancy attack (attacker's contract)

```
1 import "EtherStore.sol";
2
3 contract Attack {
4     EtherStore public etherStore;
5
6     constructor(address _etherStoreAddress) {
7         etherStore = EtherStore(_etherStoreAddress);
8     }
9     function pwnEtherStore() public payable {
10         require(msg.value >= 1 ether);
11         etherStore.depositFunds.value(1 ether)();
12         etherStore.withdrawFunds(1 ether);
13     }
14     function collectEther() public {
15         msg.sender.transfer(this.balance);
16     }
17     function () payable {
18         if (etherStore.balance > 1 ether) {
19             etherStore.withdrawFunds(1 ether);
20         }
21     }
22 }
```

# Reentrancy attack (attacker's contract)


```
1 import "EtherStore.sol";
2
3 contract Attack {
4     EtherStore public etherStore;
5
6     constructor(address _etherStoreAddress) {
7         etherStore = EtherStore(_etherStoreAddress);
8     }
9     function pwnEtherStore() public payable {
10         require(msg.value >= 1 ether);
11         etherStore.depositFunds.value(1 ether)();
12         etherStore.withdrawFunds(1 ether);
13     }
14     function collectEther() public {
15         msg.sender.transfer(this.balance);
16     }
17     function () payable {
18         if (etherStore.balance > 1 ether) {
19             etherStore.withdrawFunds(1 ether);
20         }
21     }
22 }
```

The attacker can **drain all balance** from the victim contract.

# Reentrancy attack (the fix)

```
1 contract EtherStore {
2     bool reentrancyMutex = false;
3     uint256 public withdrawLimit = 1 ether;
4     mapping(address => uint256) public lastWithdrawTime;
5     mapping(address => uint256) public balances;
6
7     function depositFunds() public payable {
8         balances[msg.sender] += msg.value;
9     }
10
11    function withdrawFunds (uint256 _weiToWithdraw) public {
12        require(balances[msg.sender] >= _weiToWithdraw);
13        require(_weiToWithdraw <= withdrawLimit);
14        require(now >= lastWithdrawTime[msg.sender] + 1 weeks);
15
16        balances[msg.sender] -= _weiToWithdraw;
17        lastWithdrawTime[msg.sender] = now;
18        reentrancyMutex = true;
19        msg.sender.transfer(_weiToWithdraw);
20        reentrancyMutex = false;
21    }
22 }
```

# Defensive driving



# Defensive programming

Like defensive driving, [defensive programming](#) requires the developer to **anticipate** what might go wrong in the software and program defensively against these anticipated issues, potentially [with the help of compiler, runtime, or even external auditors](#).

# Defensive programming

## Driving

---

Follow traffic rules  
Follow local customs

---

## Programming

---

Follow typing rules  
Follow coding conventions

---

In **normal** paradigm: expect others to follow the rules

In **defensive** paradigm: expect others to ignore / by-pass the rules

# Defensive programming

## Driving

---

Follow traffic rules  
Follow local customs

---

## Programming

---

Follow typing rules  
Follow coding conventions

---

In **normal** paradigm: expect others to follow the rules

In **defensive** paradigm: expect others to ignore / by-pass the rules

*Apply defensive actions at the cost of performance*



# Outline

- 1 Introduction: why study these bug types?
- 2 Undefined / counterintuitive behaviors
- 3 Insufficient sanitization on untrusted input
- 4 Invocation of / by untrusted logic
- 5 Hostile execution environment**
- 6 Conclusion

# Hostile environment

Sometimes the execution environment cannot be trusted as well!

# Front-running

```
1 contract FindThisHash {  
2   // the keccak-256 hash of some secret string  
3   bytes32 constant public hash  
4     = 0xb5b5b97fafd9855eec9b41f74dfb6c38f5951141f9a3ecd7f44d5479b630ee0a;  
5  
6   constructor() public payable {} // load with ether  
7  
8   function solve(string solution) public {  
9     // If you can find the pre image of the hash, receive 1000 ether  
10    require(hash == sha3(solution));  
11    msg.sender.transfer(1000 ether);  
12  }  
13 }
```

**Q:** What is the secret string?

# Front-running

```
1 contract FindThisHash {  
2   // the keccak-256 hash of some secret string  
3   bytes32 constant public hash  
4     = 0xb5b5b97fafd9855eec9b41f74dfb6c38f5951141f9a3ecd7f44d5479b630ee0a;  
5  
6   constructor() public payable {} // load with ether  
7  
8   function solve(string solution) public {  
9     // If you can find the pre image of the hash, receive 1000 ether  
10    require(hash == sha3(solution));  
11    msg.sender.transfer(1000 ether);  
12  }  
13 }
```

**Q:** What is the secret string?

**A:** Ethereum!

# Front-running

```
1 contract FindThisHash {  
2   // the keccak-256 hash of some secret string  
3   bytes32 constant public hash  
4     = 0xb5b5b97fafd9855eec9b41f74dfb6c38f5951141f9a3ecd7f44d5479b630ee0a;  
5  
6   constructor() public payable {} // load with ether  
7  
8   function solve(string solution) public {  
9     // If you can find the pre image of the hash, receive 1000 ether  
10    require(hash == sha3(solution));  
11    msg.sender.transfer(1000 ether);  
12  }  
13 }
```

**Q:** What is the secret string?

**A:** Ethereum!

A validator may see this solution, check it's validity, and then submit an equivalent transaction **with a much higher gas price** than the original transaction.

# Sandwich attack

Formal model of the automated market maker (AMM):  $x \cdot y = K$ .

# Sandwich attack

Formal model of the automated market maker (AMM):  $x \cdot y = K$ .

Example:

- Initial state:  $x_0 = 10$ ,  $y_0 = 30$ ,  $K = x_0 \cdot y_0 = 300$
- Exchange:  $x_1 = 15$ ,  $y_1 = 20$ ,  $K = x_1 \cdot y_1 = 300$ 
  - Expect  $-5$  on Token X and  $+10$  on token Y.

# Sandwich attack

Formal model of the automated market maker (AMM):  $x \cdot y = K$ .

Example:

- Initial state:  $x_0 = 10, y_0 = 30, K = x_0 \cdot y_0 = 300$
- Exchange:  $x_1 = 15, y_1 = 20, K = x_1 \cdot y_1 = 300$ 
  - Expect  $-5$  on Token X and  $+10$  on token Y.

Attack:

- Initial state:  $x_0 = 10, y_0 = 30, K = x_0 \cdot y_0 = 300$
- **Front-running**:  $x_1 = 15, y_1 = 20, K = x_1 \cdot y_1 = 300$ 
  - Attacker now holds  $-5$  Token X and  $+10$  token Y.
- Exchange:  $x_2 = 20, y_2 = 15, K = x_2 \cdot y_2 = 300$ 
  - Victim now exchanged  $-5$  Token X but only received  $+5$  token Y.
- **Back-running**:  $x_3 = 12, y_3 = 25, K = x_3 \cdot y_3 = 300$ 
  - Attacker now holds **3 Token X** and no token Y.



# Block timestamp dependence

```
1 contract Roulette {
2     uint public pastBlockTime; // Forces one bet per block
3
4     constructor() public payable {} // initially fund contract
5
6     // fallback function used to make a bet
7     function () public payable {
8         require(msg.value == 10 ether); // must send 10 ether to play
9         require(now != pastBlockTime); // only 1 transaction per block
10        pastBlockTime = now;
11        if(now % 15 == 0) { // winner
12            msg.sender.transfer(this.balance);
13        }
14    }
15 }
```

# Block timestamp dependence

```
1 contract Roulette {
2     uint public pastBlockTime; // Forces one bet per block
3
4     constructor() public payable {} // initially fund contract
5
6     // fallback function used to make a bet
7     function () public payable {
8         require(msg.value == 10 ether); // must send 10 ether to play
9         require(now != pastBlockTime); // only 1 transaction per block
10        pastBlockTime = now;
11        if(now % 15 == 0) { // winner
12            msg.sender.transfer(this.balance);
13        }
14    }
15 }
```

**The 15-second rule:** On Ethereum, a miner can post a timestamp within **15 seconds** of the block being validated. This effectively allows the miner to pre-compute an option more favorable to its chances in the lottery — **timestamps are not truly random!**

# Replay attacks

```
1 function transferProxy(  
2     address _from, address _to, uint256 _value, uint256 _fee,  
3     uint8 _v, bytes32 _r, bytes32 _s  
4 ) public returns (bool) {  
5     if (balances[_from] < _fee + _value || _fee > _fee + _value) revert();  
6  
7     uint256 nonce = nonces[_from];  
8     bytes32 h = keccak256(_from,_to,_value,_fee,nonce);  
9     if (_from != ecrecover(h,_v,_r,_s)) revert();  
10  
11     if (balances[_to] + _value < balances[_to]  
12         || balances[msg.sender] + _fee < balances[msg.sender]) revert();  
13     balances[_to] += _value;  
14     emit Transfer(_from, _to, _value);  
15  
16     balances[msg.sender] += _fee;  
17     emit Transfer(_from, msg.sender, _fee);  
18  
19     balances[_from] -= _value + _fee;  
20     nonces[_from] = nonce + 1;  
21     return true;  
22 }
```

# Replay attacks

```
1 function transferProxy(  
2     address _from, address _to, uint256 _value, uint256 _fee,  
3     uint8 _v, bytes32 _r, bytes32 _s  
4 ) public returns (bool) {  
5     if (balances[_from] < _fee + _value || _fee > _fee + _value) revert();  
6  
7     uint256 nonce = nonces[_from];  
8     bytes32 h = keccak256(_from,_to,_value,_fee,nonce);  
9     if (_from != ecrecover(h,_v,_r,_s)) revert();  
10  
11     if (balances[_to] + _value < balances[_to]  
12         || balances[msg.sender] + _fee < balances[msg.sender]) revert();  
13     balances[_to] += _value;  
14     emit Transfer(_from, _to, _value);  
15  
16     balances[msg.sender] += _fee;  
17     emit Transfer(_from, msg.sender, _fee);  
18  
19     balances[_from] -= _value + _fee;  
20     nonces[_from] = nonce + 1;  
21     return true;  
22 }
```

This function can be replayed **with another token!**

# Outline

- 1 Introduction: why study these bug types?
- 2 Undefined / counterintuitive behaviors
- 3 Insufficient sanitization on untrusted input
- 4 Invocation of / by untrusted logic
- 5 Hostile execution environment
- 6 Conclusion

# Conclusion

- Don't make assumptions about the programming language
- Don't make assumptions about user inputs
- Don't make assumptions about code you call into or code that calls into your program
- Don't make assumptions about the execution environment

〈 End 〉