

CS486/686: Introduction to Artificial Intelligence

Lecture 5 - Inference and Planning

Jesse Hoey & Victor Zhong

School of Computer Science, University of Waterloo

January 28, 2025

Readings: Poole & Mackworth Chap. 5.1-5.3 and 6.1-6.3

Problem Solving

Two methods for solving problems:

- **Procedural**

- Devise an algorithm
- Program the algorithm
- Execute the program

- **Declarative**

- Identify the knowledge needed
- Encode the knowledge in a representation (knowledge base - KB)
- Use logical consequences of KB to solve the problem

Problem Solving

Two methods for solving problems:

- **Procedural**

- “How to” knowledge
- Programs
- Meaning of symbols is meaning of computation
- Languages: C,C++,Java ...

- **Declarative**

- Descriptive knowledge
- Databases
- Meaning of symbols is meaning in world
- Languages: propositional logic, Prolog, relational databases, ...

Proof Procedures

A logic consists of

- **Syntax**: what is an acceptable sentence?
- **Semantics**: what do the sentences and symbols mean?
- **Proof procedure**: how do we construct valid proofs?

A proof: a **sequence of sentences derivable using an inference rule**

Logical Connectives

and (<u>conjunction</u>)	\wedge
or (<u>disjunction</u>)	\vee
not (<u>negation</u>)	\neg
if ... then ... (implication)	\rightarrow
... if and only if ...	\iff

Note: often logical statements with implication are written backwards: $A \rightarrow B$ is the same as $B \leftarrow A$.

Implication Truth Table

A	B	$A \rightarrow B$
F	F	T
F	T	T
T	F	F
T	T	T

(A)

(B)

If it rains, then I will carry an umbrella

Implication Truth Table

A	B	$A \rightarrow B$
F	F	T
F	T	T
T	F	F
T	T	T

(A)

(B)

If it rains, then I will carry an umbrella

If you don't study, then you will fail

Implication Truth Table

A	B	$A \rightarrow B$	$A \wedge \neg B$	$\neg(A \wedge \neg B)$	$\neg A \vee B$
F	F	T	F	T	T
F	T	T	F	T	T
T	F	F	T	F	F
T	T	T	F	T	T

(A)

(B)

no rain or I will carry an umbrella

study or you will fail

Logical Consequence

- $\{X\}$ is a set of **statements** or **premises**, made up of **propositions**
- A set of truth assignments to the propositions in $\{X\}$ is an **interpretation**
- A **model** of $\{X\}$ is an interpretation that makes $\{X\}$ true
- We say that the world in which these truth assignments hold is a **model** (a verifiable **example**) of $\{X\}$
- $\{X\}$ is **inconsistent** if it has **no model**

Logical Consequence

A statement A is a logical consequence of a set of statements $\{X\}$, if A is true in every model of $\{X\}$

If, for every set of truth assignments that hold for $\{X\}$ (for every model of $\{X\}$), some other statement (A) is always true,
Then this other statement is a **logical consequence** of $\{X\}$

Arguments and Models

P1: If I play hockey, then I'll score a goal if the goalie is not good

P2: If I play hockey, the goalie is not good

D: Therefore, if I play hockey, I'll score a goal

P: I play hockey C: I'll score a goal H: the goalie is good

$P1 : P \rightarrow (\neg H \rightarrow C)$ $P2 : P \rightarrow \neg H$ $D : P \rightarrow C$

P	C	H	$\neg H \rightarrow C$	$P1$	$P2$	D
F	F	F	F	T	T	T
F	F	T	T	T	T	T
F	T	F	T	T	T	T
F	T	T	T	T	T	T
T	F	F	F	F	T	F
T	F	T	T	T	F	F
T	T	F	T	T	T	T
T	T	T	T	T	F	T

Arguments and Models

P1: If I play hockey, then I'll score a goal if the goalie is not good

P2: If I play hockey, the goalie is not good

D: Therefore, if I play hockey, I'll score a goal

P: I play hockey C: I'll score a goal H: the goalie is good

$P1 : P \rightarrow (\neg H \rightarrow C)$ $P2 : P \rightarrow \neg H$ $D : P \rightarrow C$

P	C	H	$\neg H \rightarrow C$	$P1$	$P2$	D
F	F	F	F	T	T	T
F	F	T	T	T	T	T
F	T	F	T	T	T	T
F	T	T	T	T	T	T
T	F	F	F	F	T	F
T	F	T	T	T	F	F
T	T	F	T	T	T	T
T	T	T	T	T	F	T

Argument Validity

An argument is **valid** if any of the following is true:

- the conclusions are a **logical consequence** of the premises
- the conclusions are true in **every model** of the premises
- there is **no** situation in which the premises are all true, but the conclusions are false
- argument \rightarrow conclusions is a **tautology** (always true)

Argument Validity

An argument is **valid** if any of the following is true:

- the conclusions are a **logical consequence** of the premises
- the conclusions are true in **every model** of the premises
- there is **no** situation in which the premises are all true, but the conclusions are false
- argument \rightarrow conclusions is a **tautology** (always true)

(these four statements are identical)

Arguments and Models

P	C	H	$\neg H \rightarrow C$	$P1$	$P2$	D
F	F	F	F	T	T	T
F	F	T	T	T	T	T
F	T	F	T	T	T	T
F	T	T	T	T	T	T
T	F	F	F	F	T	F
T	F	T	T	T	F	F
T	T	F	T	T	T	T
T	T	T	T	T	F	T

- Each row is an **interpretation**: an assignment T/F to each proposition
In all the green lines, the premises are true: these interpretations are **models** of $P1$ and $P2$
- Every** model of $P1$ and $P2$ is a model of D
Therefore, D is a **logical consequence** of $P1$ and $P2$:

$$P1, P2 \models D$$

Logical Consequence

P1: Elvis is Dead

P2: Elvis is Not Dead

D: Therefore, Jerry is Alive

Is this argument valid?

Logical Consequence

P1: Elvis is Dead

P2: Elvis is Not Dead

D: Therefore, Jerry is Alive

Is this argument valid? Yes!

E: Elvis is Alive J: Jerry is Alive

E	$\neg E$	J
F	T	F
F	T	T
T	F	F
T	F	T

An argument is **valid** if there is **no** situation in which the premises are all true, but the conclusions are false

But here, there is **no model of the premises**, so the argument is **valid**

Proofs

- A **Knowledge Base** (KB) is a set of axioms
- A **proof procedure** is a way of proving theorems
- $KB \vdash g$ means g can be **derived** from KB using the proof procedure
- If $KB \vdash g$, then g is a **theorem**
- A proof procedure is **sound**:
if $KB \vdash g$ then $KB \models g$.
- A proof procedure is **complete**:
if $KB \models g$ then $KB \vdash g$.
- Two types of proof procedures:
bottom up and **top down**

Complete Knowledge

- We assume a **closed world**
 - the agent knows everything (or can prove everything)
 - if it can't prove something: must be false
 - **negation as failure**
- Another option is an **open world**:
 - the agent doesn't know everything
 - can't conclude anything from a lack of knowledge

Bottom-Up Proof

Also known as **forward chaining** - start from facts and use rules to generate all possible propositions $\text{rain} \leftarrow \text{clouds} \wedge \text{wind}$

$\text{clouds} \leftarrow \text{humid} \wedge \text{cyclone}$

$\text{clouds} \leftarrow \text{near_sea} \wedge \text{cyclone}$

$\text{wind} \leftarrow \text{cyclone}$

near_sea

cyclone

Bottom-Up Proof

Also known as **forward chaining** - start from facts and use rules to generate all possible propositions $\text{rain} \leftarrow \text{clouds} \wedge \text{wind}$

$\text{clouds} \leftarrow \text{humid} \wedge \text{cyclone}$

$\text{clouds} \leftarrow \text{near_sea} \wedge \text{cyclone}$

$\text{wind} \leftarrow \text{cyclone}$

near_sea

cyclone

$\{\text{near_sea}, \text{cyclone}\}$

Bottom-Up Proof

Also known as **forward chaining** - start from facts and use rules to generate all possible propositions $\text{rain} \leftarrow \text{clouds} \wedge \text{wind}$

$\text{clouds} \leftarrow \text{humid} \wedge \text{cyclone}$

$\text{clouds} \leftarrow \text{near_sea} \wedge \text{cyclone}$

$\text{wind} \leftarrow \text{cyclone}$

near_sea

cyclone

{near_sea, cyclone}

{near_sea, cyclone, wind}

Bottom-Up Proof

Also known as **forward chaining** - start from facts and use rules to generate all possible propositions $\text{rain} \leftarrow \text{clouds} \wedge \text{wind}$

$\text{clouds} \leftarrow \text{humid} \wedge \text{cyclone}$

$\text{clouds} \leftarrow \text{near_sea} \wedge \text{cyclone}$

$\text{wind} \leftarrow \text{cyclone}$

near_sea

cyclone

$\{\text{near_sea}, \text{cyclone}\}$

$\{\text{near_sea}, \text{cyclone}, \text{wind}\}$

$\{\text{near_sea}, \text{cyclone}, \text{wind}, \text{clouds}\}$

Bottom-Up Proof

Also known as **forward chaining** - start from facts and use rules to generate all possible propositions $\text{rain} \leftarrow \text{clouds} \wedge \text{wind}$

$\text{clouds} \leftarrow \text{humid} \wedge \text{cyclone}$

$\text{clouds} \leftarrow \text{near_sea} \wedge \text{cyclone}$

$\text{wind} \leftarrow \text{cyclone}$

near_sea

cyclone

{near_sea, cyclone}

{near_sea, cyclone, wind}

{near_sea, cyclone, wind, clouds}

{near_sea, cyclone, wind, clouds, rain}

Bottom-up proof

$C := \{\};$

repeat

 select $r \in KB$ such that

- r is $h \leftarrow b_1 \wedge \dots \wedge b_m$
- $b_i \in C \quad \forall \quad i$
- $h \notin C$

$C := C \cup \{h\}$

until no more statements can be selected

Sound and Complete

Top-Down Proof

Start from query and work backwards

`rain \leftarrow clouds \wedge wind`

`clouds \leftarrow humid \wedge cyclone`

`clouds \leftarrow near_sea \wedge cyclone`

`wind \leftarrow cyclone`

`near_sea`

`cyclone`

Start with query: if rain is proved, "yes" is the logical result (the answer to the question)

Top-Down Proof

Start from query and work backwards

```
rain  $\leftarrow$  clouds  $\wedge$  wind  
clouds  $\leftarrow$  humid  $\wedge$  cyclone  
clouds  $\leftarrow$  near_sea  $\wedge$  cyclone  
wind  $\leftarrow$  cyclone  
near_sea  
cyclone
```

Start with query: if rain is proved, "yes" is the logical result (the answer to the question)

```
yes  $\leftarrow$  rain
```

Top-Down Proof

Start from query and work backwards

```
rain  $\leftarrow$  clouds  $\wedge$  wind  
clouds  $\leftarrow$  humid  $\wedge$  cyclone  
clouds  $\leftarrow$  near_sea  $\wedge$  cyclone  
wind  $\leftarrow$  cyclone  
near_sea  
cyclone
```

```
yes  $\leftarrow$  rain  
yes  $\leftarrow$  clouds  $\wedge$  wind
```

Top-Down Proof

Start from query and work backwards

```
rain  $\leftarrow$  clouds  $\wedge$  wind  
clouds  $\leftarrow$  humid  $\wedge$  cyclone  
clouds  $\leftarrow$  near_sea  $\wedge$  cyclone  
wind  $\leftarrow$  cyclone  
near_sea  
cyclone
```

```
yes  $\leftarrow$  rain  
yes  $\leftarrow$  clouds  $\wedge$  wind  
yes  $\leftarrow$  near_sea  $\wedge$  cyclone  $\wedge$  wind
```

Top-Down Proof

Start from query and work backwards

`rain \leftarrow clouds \wedge wind`

`clouds \leftarrow humid \wedge cyclone`

`clouds \leftarrow near_sea \wedge cyclone`

`wind \leftarrow cyclone`

`near_sea`

`cyclone`

`yes \leftarrow rain`

`yes \leftarrow clouds \wedge wind`

`yes \leftarrow near_sea \wedge cyclone \wedge wind`

`yes \leftarrow near_sea \wedge cyclone \wedge cyclone`

Top-Down Proof

Start from query and work backwards

```
rain ← clouds ∧ wind
clouds ← humid ∧ cyclone
clouds ← near_sea ∧ cyclone
wind ← cyclone
near_sea
cyclone
```

```
yes ← rain
yes ← clouds ∧ wind
yes ← near_sea ∧ cyclone ∧ wind
yes ← near_sea ∧ cyclone ∧ cyclone
yes ← near_sea ∧ cyclone
```

Top-Down Proof

Start from query and work backwards

```
rain  $\leftarrow$  clouds  $\wedge$  wind  
clouds  $\leftarrow$  humid  $\wedge$  cyclone  
clouds  $\leftarrow$  near_sea  $\wedge$  cyclone  
wind  $\leftarrow$  cyclone  
near_sea  
cyclone
```

```
yes  $\leftarrow$  rain  
yes  $\leftarrow$  clouds  $\wedge$  wind  
yes  $\leftarrow$  near_sea  $\wedge$  cyclone  $\wedge$  wind  
yes  $\leftarrow$  near_sea  $\wedge$  cyclone  $\wedge$  cyclone  
yes  $\leftarrow$  near_sea  $\wedge$  cyclone  
yes  $\leftarrow$  cyclone
```


Top-Down Proof

Start from query and work backwards

```
rain  $\leftarrow$  clouds  $\wedge$  wind  
clouds  $\leftarrow$  humid  $\wedge$  cyclone  
clouds  $\leftarrow$  near_sea  $\wedge$  cyclone  
wind  $\leftarrow$  cyclone  
near_sea  
cyclone
```

```
yes  $\leftarrow$  rain  
yes  $\leftarrow$  clouds  $\wedge$  wind  
yes  $\leftarrow$  near_sea  $\wedge$  cyclone  $\wedge$  wind  
yes  $\leftarrow$  near_sea  $\wedge$  cyclone  $\wedge$  cyclone  
yes  $\leftarrow$  near_sea  $\wedge$  cyclone  
yes  $\leftarrow$  cyclone  
yes  $\leftarrow$ 
```

Top-Down Interpreter

solve($q_1 \wedge \dots \wedge q_k$):

$ac := \text{"yes"} \leftarrow q_1 \wedge \dots \wedge q_k''$

repeat

select a conjunct q_i from body of ac

choose a statement C from KB with q_i as head

replace q_i in body of ac by body of C

until ac is an answer

Beyond propositions: Individuals and Relations

- KB can contain **relations** : `part_of(C,A)` is true if C is a “part of” A (in the world)
- KB can contain **quantification** : `part_of(C,A)` holds $\forall C, A$
- Proof procedure is the same, with a few extra bits to handle relations & quantification


Planning

- Planning is **deciding what to do** based on an agent's ability, its goals, and the state of the world
- Planning is finding a **sequence of actions to solve a goal**
- Initial assumptions:
 - A **single** agent
 - The world is **deterministic**
 - There are **no exogenous events** outside of the control of the agent that change the state of the world
 - The agent knows what state it is in (full **observability**)
 - **Time progresses discretely** from one state to the next
 - **Goals are predicates** of states that need to be achieved or maintained (no complex goals)

Actions

- A deterministic **action** is a **partial function from states to states**
- **partial** function: some actions not possible in some states
- The **preconditions** of an action specify when the action can be carried out
- The **effect** of an action specifies the resulting state

Delivery Robot Example



Features (Variables):

RLoc – Rob's location

(4-valued: {cs,off,mr,lab})

RHC – Rob has coffee (binary)

SWC – Sam wants coffee (binary)

MW – Mail is waiting (binary)

RHM – Rob has mail (binary)

Actions:

mc – move clockwise

mcc – move counterclockwise

puc – pickup coffee

dc – deliver coffee

pum – pickup mail

dm – deliver mail

Explicit State-Space Representation

State	Action	Resulting State
$\langle lab, \neg rhc, swc, \neg mw, rhm \rangle$	<i>mc</i>	$\langle mr, \neg rhc, swc, \neg mw, rhm \rangle$
$\langle lab, \neg rhc, swc, \neg mw, rhm \rangle$	<i>mcc</i>	$\langle off, \neg rhc, swc, \neg mw, rhm \rangle$
$\langle off, \neg rhc, swc, \neg mw, rhm \rangle$	<i>dm</i>	$\langle off, \neg rhc, swc, \neg mw, \neg rhm \rangle$
$\langle off, \neg rhc, swc, \neg mw, rhm \rangle$	<i>mcc</i>	$\langle cs, \neg rhc, swc, \neg mw, rhm \rangle$
$\langle off, \neg rhc, swc, \neg mw, rhm \rangle$	<i>mc</i>	$\langle lab, \neg rhc, swc, \neg mw, rhm \rangle$
...

Feature-Based Representation of Actions

For each action:

- **precondition** is a proposition that specifies when the action can be carried out.

For each feature:

- **causal rules** that specify when the feature gets a new value and
- **frame rules** that specify when the feature keeps its value.

Notation:

- Features are capitalized (e.g. *Rloc*, *RHC*)
- Values of the features are not (e.g. *Rloc* = *cs*, *rhc*, $\neg rhc$)
- If *X* is a feature, then *X'* is the feature after an action is carried out

Example feature-based representation

Precondition of pick-up coffee (*puc*):

$$RLoc = cs \wedge \neg rhc$$

Rules for location is *cs* (specifies *RLoc'*):

$$RLoc' = cs \leftarrow RLoc = off \wedge Act = mcc$$

$$RLoc' = cs \leftarrow RLoc = mr \wedge Act = mc$$

$$RLoc' = cs \leftarrow RLoc = cs \wedge Act \neq mcc \wedge Act \neq mc$$

Rules for “robot has coffee” (specifies *rhc'*):

$$\text{(frame rule)} \quad RHC' = true \leftarrow RCH = true \wedge Act \neq dc$$

$$\text{(or } rhc' \leftarrow rhc \wedge Act \neq dc)$$

$$\text{(causal rule)} \quad RHC' = true \leftarrow Act = puc \text{ (or } rhc' \leftarrow Act = puc)$$

Planning

Given:

- A description of the effects and preconditions of the actions
- A description of the initial state
- A goal to achieve

Find a sequence of actions that is possible and will result in a state satisfying the goal

Forward Planning

Idea: search in the state-space graph

- The nodes represent the states
- The arcs correspond to the actions: The arcs from a state s represent all of the actions that are legal in state s
- A plan is a path from the state representing the initial state to a state that satisfies the goal.
- Can use any of the search techniques from Lecture 3
- **heuristics** important
- A tutorial by Malte Helmert on Heuristics for Deterministic Planning:
https://ai.dmi.unibas.ch/misc/tutorial_aaai2015/

Example State-Space Graph

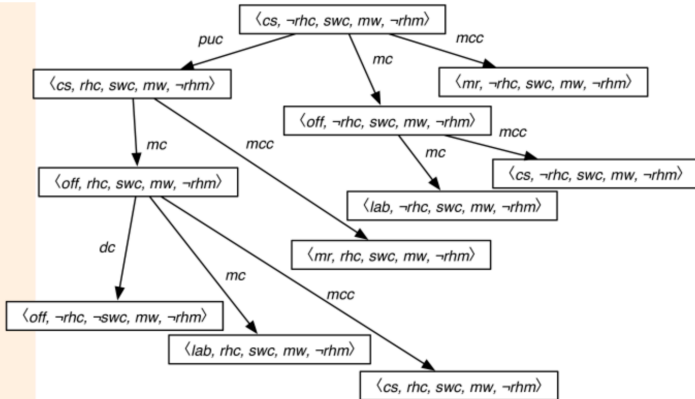


Figure 6.2: Part of the search space for a state-space planner

Regression Planning

Idea: search **backwards** from the goal description: nodes correspond to subgoals, and arcs to actions

- Nodes are propositions: a formula made up of assignments of values to features
- Arcs correspond to actions that can achieve one of the goals
- Neighbors of a node N associated with arc A specify what must be true immediately before A so that N is true immediately after
- The start node is the goal to be achieved
- $\text{goal}(N)$ is true if N is a proposition that is true of the initial state

Next

- Supervised learning (Poole & Mackworth chapter 7.1-7.3.1, 7.4-7.4.1)