

# CS486/686: Introduction to Artificial Intelligence

## Lecture 9c - Reinforcement Learning

Jesse Hoey & Victor Zhong

School of Computer Science, University of Waterloo

March 27, 2025

Readings: Poole & Mackworth Chap. 13.1-13.9

# Reinforcement Learning

What should an agent do given:

- **Prior knowledge:** possible states of the world  
possible actions
- **Observations:** current state of world  
immediate reward / punishment
- **Goal:** act to maximize accumulated reward

Like decision-theoretic planning, except model of dynamics and model of reward not given.

# Experiences

- We assume there is a **sequence of experiences**:

*state, action, reward, state, action, reward, ....*

- What should the agent do next?
- It must decide whether to:
  - **explore** to gain more knowledge
  - **exploit** the knowledge it has already discovered

## Reinforcement Learning: “Bandit” problem



Each machine has a  $Pr(win)$  ... but you don't know what it is...

**Which machine should you play?**

# Why is reinforcement learning hard?

- **credit assignment** problem: what actions are responsible for the reward may have occurred a **long time before** the reward was received.
- The **long-term effect** of an action of the robot depends on what it will do in the future, but it does not have a model to predict this.
- The **explore-exploit dilemma**: at each time should the robot be greedy or inquisitive?

# Reinforcement learning: main approaches

- **Model-Free RL** learn  $Q^*(s, a)$ , use this to guide action.
- **Model Based RL**: **learn a model** consisting of state transition function  $P(s'|a, s)$  and reward function  $R(s, a, s')$ ; **solve** this as an MDP.
- search through a space of policies

# Temporal Differences

- Suppose we have a sequence of values:

$$v_1, v_2, v_3, \dots$$

And want a **running estimate** of the average of the first  $k$  values:

$$A_k = \frac{v_1 + \dots + v_k}{k}$$

## Temporal Differences (cont)

- When a new value  $v_k$  arrives:

$$\begin{aligned}A_k &= \frac{v_1 + \cdots + v_{k-1} + v_k}{k} \\kA_k &= v_1 + \cdots + v_{k-1} + v_k \\&= (k-1)A_{k-1} + v_k \\A_k &= \frac{k-1}{k}A_{k-1} + \frac{1}{k}v_k\end{aligned}$$

Let  $\alpha = \frac{1}{k}$ , then

$$\begin{aligned}A_k &= (1 - \alpha)A_{k-1} + \alpha v_k \\&= A_{k-1} + \alpha(v_k - A_{k-1})\end{aligned}$$

### “TD formula”

- Often we use this update with  $\alpha$  **fixed**.



# Q-learning

- **Idea:** store  $Q[\text{State}, \text{Action}]$ ; update this as in **asynchronous value iteration**, but using experience (empirical probabilities and rewards).
- Suppose the agent has an experience  $\langle s, a, r, s' \rangle$
- This provides one piece of data to update  $Q[s, a]$ .
- The experience  $\langle s, a, r, s' \rangle$  provides the data point:

$$r + \gamma \max_{a'} Q[s', a']$$

which can be used in the **TD formula** giving:

$$Q[s, a] \leftarrow Q[s, a] + \alpha \left( r + \gamma \max_{a'} Q[s', a'] - Q[s, a] \right)$$

# Q-learning

initialize  $Q[S, A]$  arbitrarily

observe current state  $s$

**repeat forever:**

    select and carry out an action  $a$

    observe reward  $r$  and state  $s'$

$Q[s, a] \leftarrow Q[s, a] + \alpha (r + \gamma \max_{a'} Q[s', a'] - Q[s, a])$

$s \leftarrow s'$ ;

**end repeat**

# Properties of Q-learning

- Q-learning **converges to the optimal policy**, no matter what the agent does, as long as it **tries each action in each state enough (infinitely often)**.
- But what should the agent do?
  - **exploit**: when in state  $s$ , select the action that maximizes  $Q[s, a]$
  - **explore**: select another action

# Exploration Strategies

- The  **$\epsilon$ -greedy** strategy: choose a random action with probability  $\epsilon$  and choose a best action with probability  $1 - \epsilon$ .
- **Softmax** action selection: in state  $s$ , choose action  $a$  with probability

$$\frac{e^{Q[s,a]/\tau}}{\sum_a e^{Q[s,a]/\tau}}$$

where  $\tau > 0$  is the **temperature**. Good actions are chosen more often than bad actions;  $\tau$  defines how often good actions are chosen. For  $\tau \rightarrow \infty$ , all actions are equiprobable. For  $\tau \rightarrow 0$ , only the best is chosen.


# Exploration Strategies

- **optimism in the face of uncertainty**: initialize  $Q$  to values that encourage exploration.
- **Upper Confidence Bound (UCB)**: Also store  $N[s, a]$  (number of times that state-action pair has been tried) and use

$$\arg \max_a \left[ Q(s, a) + k \sqrt{\frac{N[s]}{N[s, a]}} \right]$$

where  $N[s] = \sum_a N[s, a]$


## Example: Studentbot



state variables:

- **tired**: studentbot is tired (no/a bit/very)
- **passtest**: studentbot passes test (no/yes)
- **knows**: studentbot's state of knowledge (nothing/a bit/a lot/everything)
- **goodtime**: studentbot has a good time (no/yes)


## Example: Studentbot



studentbot actions:

- **study**: studentbot's knowledge increases, studentbot gets tired
- **sleep**: studentbot gets less tired
- **party**: studentbot has a good time, but gets tired and loses knowledge
- **take test**: studentbot takes a test (can take test anytime)

## Example: Studentbot



studentbot rewards:

- **+20** if studentbot passes the test
- **+2** if studentbot has a good time when not very tired

basic tradeoff: short term vs. long-term rewards



## Studentbot Policy

| tired | know       | policy    | value | policy    | value |
|-------|------------|-----------|-------|-----------|-------|
| no    | nothing    | party     | 10.53 | study     | 16.7  |
| no    | a bit      | party     | 11.96 | study     | 20.8  |
| no    | a lot      | sleep     | 18.05 | study     | 25.7  |
| no    | everything | take test | 33.73 | take test | 31.6  |
| a bit | nothing    | sleep     | 9.47  | study     | 15.97 |
| a bit | a bit      | take test | 10.02 | study     | 19.94 |
| a bit | a lot      | sleep     | 27.54 | study     | 26.0  |
| a bit | everything | take test | 32.39 | take test | 31.6  |
| very  | nothing    | sleep     | 9.52  | sleep     | 15.1  |
| very  | a bit      | sleep     | 11.08 | sleep     | 18.7  |
| very  | a lot      | sleep     | 27.52 | sleep     | 23.1  |
| very  | everything | take test | 20.47 | sleep     | 28.4  |

Policy after 100,000 RL steps for  
pass=no, gt=no

Optimal Value and  
Policy

## Studentbot Policy

| tired | know       | policy    | value | policy    | value |
|-------|------------|-----------|-------|-----------|-------|
| no    | nothing    | party     | 10.53 | study     | 16.7  |
| no    | a bit      | study     | 14.61 | study     | 20.8  |
| no    | a lot      | study     | 25.3  | study     | 25.7  |
| no    | everything | take test | 29.9  | take test | 31.6  |
| a bit | nothing    | take test | 9.52  | study     | 15.97 |
| a bit | a bit      | take test | 16.06 | study     | 19.94 |
| a bit | a lot      | sleep     | 21.43 | study     | 26.0  |
| a bit | everything | sleep     | 39.8  | take test | 31.6  |
| very  | nothing    | sleep     | 9.47  | sleep     | 15.1  |
| very  | a bit      | sleep     | 13.41 | sleep     | 18.7  |
| very  | a lot      | sleep     | 17.8  | sleep     | 23.1  |
| very  | everything | sleep     | 28.32 | sleep     | 28.4  |

Policy after 200,000 RL steps for  
pass=no, gt=no

Optimal Value and  
Policy

# Off/On-policy Learning

- Q-learning does **off-policy** learning: it learns the value of the optimal policy, no matter what it does.
- This could be bad if the exploration policy is dangerous.
- **On-policy** learning learns the value of the policy being followed.  
e.g., act greedily 80% of the time and act randomly 20% of the time
- If the agent is actually going to explore, it may be better to optimize the actual policy it is going to do.
- **SARSA** uses the experience  $\langle s, a, r, s', a' \rangle$  to update  $Q[s, a]$ .
- **SARSA** is **on-policy** because it uses empirical values for  $s'$  and  $a'$

# SARSA

initialize  $Q[S, A]$  arbitrarily

observe current state  $s$

select action  $a$  using a policy based on  $Q$  **includes exploration**

**repeat forever:**

    carry out an action  $a$

    observe reward  $r$  and state  $s'$

    select action  $a'$  using a policy based on  $Q$

$$Q[s, a] \leftarrow Q[s, a] + \alpha (r + \gamma Q[s', a'] - Q[s, a])$$

$s \leftarrow s'$ ;

$a \leftarrow a'$ ;

**end-repeat**

# Model-based Reinforcement Learning

- **Model-based** reinforcement learning uses the experiences in a more effective manner.
- It is used when **collecting experiences is expensive** (e.g., in a robot or an online game), and you can do lots of computation between each experience.
- **Idea**: learn the MDP and **interleave acting and planning**.
- After each experience, update probabilities and the reward, then do some steps of **asynchronous value iteration**.

## Model-based learner

Data Structures:  $Q[S, A]$ ,  $T[S, A, S]$ ,  $R[S, A]$

Assign  $Q$ ,  $R$  arbitrarily,  $T$  = prior counts

$\alpha$  is learning rate

observe current state  $s$

**repeat forever:**

select and carry out action  $a$

observe reward  $r$  and state  $s'$

$$T[s, a, s'] \leftarrow T[s, a, s'] + 1$$

$$R[s, a] \leftarrow \alpha \times r + (1 - \alpha) \times R[s, a]$$

**repeat for a while (asynchronous VI):**

select state  $s_1$ , action  $a_1$

$$\text{let } P = \sum_{s_2} T[s_1, a_1, s_2]$$

$$Q[s_1, a_1] \leftarrow \sum_{s_2} \frac{T[s_1, a_1, s_2]}{P} (R[s_1, a_1] + \gamma \max_{a_2} Q[s_2, a_2])$$


**end-repeat**

$$s \leftarrow s';$$


**end-repeat**

# Large State Spaces

- **Computer Go:**  $3^{361}$  states



- **Atari Games**  $210 \times 160 \times 3$  dimensions (pixels)



# Q-function Approximations

- Let  $s = (x_1, x_2, \dots, x_N)^T$  where  $x$  are **features**
- **Linear**

$$Q_{\mathbf{w}}(s, a) \approx \sum_i w_{ai} x_i$$

- **Non-linear** (e.g. neural network  $g$ )

$$Q_{\mathbf{w}}(s, a) \approx g(\mathbf{x}; \mathbf{w}_a)$$



## Recall: Logistic Regression

Logistic function of linear weighted inputs:

$$\hat{Y}^{\bar{w}}(e) = f(w_0 + w_1 X_1(e) + \cdots + w_n X_n(e)) = f\left(\sum_{i=0}^n w_i X_i(e)\right)$$

The sum of squares error is:

$$Error(E, \bar{w}) = \sum_{e \in E} \left[ Y(e) - f\left(\sum_{i=0}^n w_i X_i(e)\right) \right]^2$$

The partial derivative with respect to weight  $w_i$  is:

$$\frac{\partial Error(E, \bar{w})}{\partial w_i} = -2\delta f' \left( \sum_i w_i X_i(e) \right) X_i(e)$$

where  $\delta = (Y(e) - f(\sum_{i=0}^n w_i X_i(e)))$ .

Thus, each example  $e$  updates each weight  $w_i$  by

$$w_i \leftarrow w_i + \eta \delta f' \left( \sum_i w_i X_i(e) \right) X_i(e)$$

## Approximating the Q-function

- assign weights  $\mathbf{w} = \bar{\mathbf{w}} = \langle w_0, \dots, w_n \rangle$  arbitrarily
- for experience tuple  $s, a, r, s', a'$  we have:
  - target Q-function:  $[R(s) + \gamma \max_a Q_{\bar{\mathbf{w}}}(s', a)]$  or  $[R(s) + \gamma Q_{\bar{\mathbf{w}}}(s', a')]$
  - current Q-function:  $Q_{\mathbf{w}}(s, a)$
- Squared error:

$$Err(\mathbf{w}) = \frac{1}{2} \left[ Q_{\mathbf{w}}(s, a) - R(s) - \gamma \max_{a'} Q_{\bar{\mathbf{w}}}(s', a') \right]^2$$

- Gradient:

$$\frac{\partial Err}{\partial \mathbf{w}} = \left[ Q_{\mathbf{w}}(s, a) - R(s) - \gamma \max_{a'} Q_{\bar{\mathbf{w}}}(s', a') \right] \frac{\partial Q_{\mathbf{w}}(s, a)}{\partial \mathbf{w}}$$

- Weight update:

$$\mathbf{w} \leftarrow \mathbf{w} - \alpha \frac{\partial Err}{\partial \mathbf{w}}$$

- Occasionally:

## SARSA with linear function approximation

Given  $\gamma$ :discount factor;  $\alpha$ :learning rate

Assign weights  $\bar{\mathbf{w}} = \langle w_0, \dots, w_n \rangle$  arbitrarily

**begin**

observe current state  $s$

select action  $a$

**repeat forever:**

carry out action  $a$

observe reward  $r$  and state  $s'$

select action  $a'$  (using a policy based on  $Q_{\bar{\mathbf{w}}}$ )

let  $\delta = r + \gamma Q_{\bar{\mathbf{w}}}(s', a') - Q_{\bar{\mathbf{w}}}(s, a)$

For  $i = 0$  to  $n$

$$w_i \leftarrow w_i + \alpha \times \delta \times \frac{\partial Q_{\bar{\mathbf{w}}}(s, a)}{\partial w_i}$$

$s \leftarrow s'$ ;  $a \leftarrow a'$ ; (sometimes:  $\bar{\mathbf{w}} \leftarrow \mathbf{w}$ ;) )

**end-repeat**

**end**

# Convergence

- Linear Q-learning ( $Q_{\mathbf{w}}(s, a) \approx \sum_i w_{ai}x_i$ ) converges under same conditions as Q-learning

$$w_i \leftarrow w_i + \alpha [Q_{\mathbf{w}}(s, a) - R(s) - \gamma Q_{\mathbf{w}}(s', a')] x_i$$

- Non-linear Q-learning (e.g. neural network,  $Q_{\mathbf{w}}(s, a) \approx g(\mathbf{x}; \mathbf{w})$ ) may diverge
- Adjusting  $\mathbf{w}$  to increase  $Q$  at  $(s, a)$  might introduce errors at nearby state-action pairs.

# Mitigating Divergence

Two tricks used in practice:

1. **Experience Replay**
2. **Use two  $Q$  function** (two networks):
  - $Q$  network (currently being updated)
  - Target network (occasionally updated)

# Experience Replay

- **Idea:** Store previous experiences  $(s, a, r, s', a')$  in a buffer and sample a mini-batch of previous experiences at each step to learn by Q-learning
- **Breaks correlations** between successive updates (more stable learning)
- Few interactions with environment needed to converge (greater **data efficiency**)

# Target Network

- **Idea**: use a separate target network that is **updated only periodically**
- target network has weights  $\bar{\mathbf{w}}$  and computes  $Q_{\bar{\mathbf{w}}}(s, a)$
- repeat for each  $(s, a, r, s', a')$  in **mini-batch**:

$$\mathbf{w} \leftarrow \mathbf{w} + \alpha \left[ Q_{\mathbf{w}}(s, a) - R(s) - \gamma Q_{\bar{\mathbf{w}}}(s', a') \right] \frac{\partial Q_{\mathbf{w}}(s, a)}{\partial \mathbf{w}}$$

- $\bar{\mathbf{w}} \leftarrow \mathbf{w}$

## Deep Q-Network

Assign weights  $\bar{\mathbf{w}} = \langle w_0, \dots, w_n \rangle$  at random in  $[-1, 1]$

**begin**

observe current state  $s$

select action  $a$

**repeat forever:**

carry out action  $a$

observe reward  $r$  and state  $s'$

select action  $a'$  (using a policy based on  $Q_{\bar{\mathbf{w}}}$ )

add  $(s, a, r, s', a')$  to experience buffer

Sample mini-batch of experiences from buffer

For each experience  $(\hat{s}, \hat{a}, \hat{r}, \hat{s}', \hat{a}')$  in mini-batch:

let  $\delta = \hat{r} + \gamma Q_{\bar{\mathbf{w}}}(\hat{s}', \hat{a}') - Q_{\mathbf{w}}(\hat{s}, \hat{a})$

$\mathbf{w} \leftarrow \mathbf{w} + \alpha \times \delta \times \frac{\partial Q_{\mathbf{w}}(\hat{s}, \hat{a})}{\partial \mathbf{w}}$

$s \leftarrow s'; a \leftarrow a';$


every  $c$  steps, update target  $\bar{\mathbf{w}} \leftarrow \mathbf{w}$

**end-repeat**

**end**




# Deep Q-Network for Atari



from: Mnih et. al. Human-level control through deep reinforcement learning.  
Nature 18(7540):529–533 2015.

# Deep Q-Network vs. Linear Approx.



## Policy search using policy gradient

- Instead of approximating  $Q$ , we can also approximate policy  $\pi$
- $\pi(a|s)$  is now a function that maps state  $s$  to a **distribution** over actions  $a$
- Let  $s = (x_1, x_2, \dots, x_N)^T$  where  $x$  are **features**
- **Linear**

$$Q_{\mathbf{w}}(s, a) \approx \sum_i w_{ai} x_i$$

$$\pi_{\theta}(a|s) \approx \sum_i \theta_{ai} x_i$$

- **Non-linear** (e.g. neural network  $g$ )

$$Q_{\mathbf{w}}(s, a) \approx g(\mathbf{x}; \mathbf{w}_a)$$

$$\pi_{\theta}(a|s) \approx g(\mathbf{x}; \theta_a)$$

# Policy search using policy gradient

- How do we **change parameters  $\theta$  of policy  $\pi_\theta$**  to maximize objective?
- What is objective?
- $V, Q$ : be value functions under  $\pi_\theta$
- $d^\pi(s) = \lim_{t \rightarrow \infty} P(s_t = s | s_0, \pi_\theta)$ : be **stationary distribution of Markov chain** for  $\pi_\theta$ 
  - Imagine traveling Markov chain's states forever, eventually probability of ending up with one state becomes unchanged
- For more see:  
[lilianweng.github.io/posts/2018-04-08-policy-gradient](https://lilianweng.github.io/posts/2018-04-08-policy-gradient)
- Want to maximize the objective:

$$J(\theta) = \sum_{s \in S} d^\pi(s) V(s) = \sum_{s \in S} d^\pi(s) \sum_{a \in A} \pi_\theta(a|s) Q(s, a)$$

## How to find a good policy?

- Gradient is tricky because depends on both action selection  $\pi_\theta$  and state distribution  $d(s)$
- But updating policy changes state distribution!
- Policy gradient theorem (Sutton & Barto 2017; optional) allows us to simplify gradient

$$\begin{aligned}\nabla_\theta J(\theta) &= \nabla_\theta \sum_{s \in \mathcal{S}} d^\pi(s) \sum_{a \in \mathcal{A}} \pi_\theta(a|s) Q(s, a) \\ &\approx \sum_{s \in \mathcal{S}} d^\pi(s) \sum_{a \in \mathcal{A}} Q(s, a) \nabla_\theta \pi_\theta(a|s) \\ &= \sum_{s \in \mathcal{S}, a \in \mathcal{A}} d^\pi(s) \pi_\theta(a|s) Q(s, a) \frac{\nabla_\theta \pi_\theta(a|s)}{\pi_\theta(a|s)} \\ &= \mathbb{E}_\pi [Q(s, a) \nabla_\theta \ln \pi_\theta(a|s)]\end{aligned}$$

where  $\mathbb{E}_\pi$  refers to  $\mathbb{E}_{s \sim d^\pi, a \sim \pi_\theta}$  where both state and action distributions follow the policy (on policy).

# REINFORCE algorithm (Williams 1992)

Use sampled trajectories to compute expectation over  $Q(s, a)$

$$\begin{aligned}\nabla_{\theta} J(\theta) &\approx \mathbb{E}_{\pi} [Q(s, a) \nabla_{\theta} \ln \pi_{\theta}(a|s)] \\ &= \mathbb{E}_{\pi} \left[ \left( \sum_{k=0}^{\infty} \gamma^k r_{t+k+1} \right) \nabla_{\theta} \ln \pi_{\theta}(a_t | s_t) \right]\end{aligned}$$

Initialize  $\theta$  to random

Sample one trajectory using  $\pi_{\theta}$ :  $s_1, a_1, r_2, s_2, a_2 \dots s_T$

**For**  $t = 1, 2 \dots T$

    Estimate  $G_t = \sum_{k=0}^{T-t-2} \gamma^k r_{t+k+1}$

    Update  $\theta \leftarrow \theta + \alpha \gamma^t G_t \nabla_{\theta} \ln \pi_{\theta}(a_t | s_t)$

**end**

# REINFORCE algorithm (Williams 1992)

- REINFORCE sampling trajectories to estimate gradient results in high variance
- Idea: subtract a baseline value from return  $G_t$  to reduce variance of gradient estimation
- Use **Advantage**  $A(s_t, a_t) = Q(s_t, a_t) - V(s_t)$  instead of  $G_t$

# Actor Critic algorithm (Degrís, White, & Sutton 2012)

- AC learns **two models, Actor  $\pi_\theta$  and Critic  $Q_w$**
- $\alpha$  and  $\beta$  are learning rates
- Use **Actor** to update  $\pi_\theta(a|s)$  in the direction suggested by critic

Initialize  $s$ ,  $w$ ,  $\theta$  to random

Sample  $a \sim \pi_\theta(a|s)$

**For**  $t = 1, 2 \dots T$

Sample reward  $r_t \sim R(s, a)$ , next state  $s' \sim P(s'|s, a)$

Sample next action  $a' \sim \pi_\theta(a', s')$

Update  $\theta \leftarrow \theta + \alpha Q_w(s, a) \nabla_\theta \ln \pi_\theta(a|s)$

Compute TD correction  $\delta_t = r_t + \gamma Q_w(s', a') - Q_w(s, a)$

Update  $w \leftarrow w + \beta \delta_t \nabla_w Q_w$

Update  $a \leftarrow a'$ ,  $s \leftarrow s'$

**end**



# Bayesian Reinforcement Learning

- **Include the parameters** (transition function and observation function) in the state space
- **Model-based learning** through inference (belief state)
- State space is now continuous, **belief space is a space of continuous functions**
- Can mitigate complexity by **modeling reachable beliefs**
- **optimal** exploration-exploitation tradeoff.

## Next:

- Recap of class + final review
- Research highlight
  - Jesse: Active inference and free energy
  - Victor: History of natural language processing