

CS486/686: Introduction to Artificial Intelligence

Lecture 1 - Introduction

Jesse Hoey & Victor Zhong

School of Computer Science, University of Waterloo

January 7, 2025

Readings: Poole & Mackworth 1.1

People, Books & Web

Instructors: Jesse Hoey (S3 & S2) & Victor Zhong (S1)

TAs:

- Negar Arabzadehghahyazi (narabzad)
- Arthur Chen (h559chen)
- Varun Dhanraj (vdhanraj)
- Yuxuan Li (y624li)
- Tien Dat Nguyen (tdnguyen)
- Ruoxi Ning (r2ning)
- Hala Sheta (hsheta)
- Shivani Upadhyay (sjupadhy)
- Dake Zhang (d346zhan)

Lectures:

- Section 3: T/Th 1:00pm-2:20pm in MC 2038 (Jesse)
- Section 2: T/Th 8:30am-9:50am in MC 2038 (Jesse)
- Section 1: T/Th 2:30-3:50pm in MC 2017 (Victor)

Office hours:

Wednesday 2:00-3:00pm DC 3613/online (Jesse)

Tuesday 10:30-11:30am DC 2637 (Victor)

Office hours (TA): Notified near assignment due dates

Grade Breakdown

CS486 (Undergraduate Students)

- 4 Assignments (40%: 10% each) – approx deadlines: weeks 5, 7, 9, 12
- 1 midterm exam (15%) – Feb 13th 7:00 pm to 8:50 pm
- 1 final exam (45%) (must pass to pass the course)
- Optional project (5% bonus, proposal due at midterm)

CS686 (Graduate Students)

- 4 Assignments (25%: 6.25% each) – same deadlines
- 1 midterm exam (10%) – same time
- 1 final exam (35%)
- 1 course project report (30%, proposal due at midterm)

Grade Breakdown

CS486 (Undergraduate Students)

- 4 Assignments (40%: 10% each) – approx deadlines: weeks 5, 7, 9, 12
- 1 midterm exam (15%) – Feb 13th 7:00 pm to 8:50 pm
- 1 final exam (45%) (must pass to pass the course)
- Optional project (5% bonus, proposal due at midterm)

CS686 (Graduate Students)

- 4 Assignments (25%: 6.25% each) – same deadlines
- 1 midterm exam (10%) – same time
- 1 final exam (35%)
- 1 course project report (30%, proposal due at midterm)

Students wishing to write a project (and all CS686 students) **must** submit a project proposal

Grade Breakdown

CS486 (Undergraduate Students)

- 4 Assignments (40%: 10% each) – approx deadlines: weeks 5, 7, 9, 12
- 1 midterm exam (15%) – Feb 13th 7:00 pm to 8:50 pm
- 1 final exam (45%) (must pass to pass the course)
- Optional project (5% bonus, proposal due at midterm)

CS686 (Graduate Students)

- 4 Assignments (25%: 6.25% each) – same deadlines
- 1 midterm exam (10%) – same time
- 1 final exam (35%)
- 1 course project report (30%, proposal due at midterm)

Students wishing to write a project (and all CS686 students) **must** submit a project proposal

The final is cumulative (i.e., covering all course material) with a focus on the post-midterm material

Assignment Deadlines

- Assignment deadline times are at 11:59pm waterloo time on the day indicated.
- All assignment deadlines are fixed and will not be extended.
- Every assignment has a **grace period** of 72 hours.
- If you use university approved STA/VIF on assignment, you enter grace period: no late penalty.
- No assignment will be accepted after the grace period ends.
- No assignment questions will be answered during the grace period.
- Any student missing the midterm with university approved STA/VIF will have the midterm weight shifted to the final.

Course Project

Optional for CS486 students (5% bonus)

Mandatory for CS686 students (30% of grade)

You **must** submit a correctly constructed and formatted proposal by the midterm – will be pass/fail with **no mark**

Final project dues before the final exam

Course Project

Optional for CS486 students (5% bonus)

Mandatory for CS686 students (30% of grade)

You **must** submit a correctly constructed and formatted proposal by the midterm – will be pass/fail with **no mark**

Final project dues before the final exam

Individual project (CS686)

Group project (CS486, up to 3 members):

- **Must be substantially** more involved than individual projects
- Each team member's contribution **to the project must be clearly and specifically described in both the proposal and report**
- There must be more papers referenced and discussed for team projects

Course Project

Optional for CS486 students (5% bonus)

Mandatory for CS686 students (30% of grade)

You **must** submit a correctly constructed and formatted proposal by the midterm – will be pass/fail with **no mark**

Final project dues before the final exam

Individual project (CS686)

Group project (CS486, up to 3 members):

- **Must be substantially** more involved than individual projects
- Each team member's contribution **to the project must be clearly and specifically described in both the proposal and report**
- There must be more papers referenced and discussed for team projects

https:

[//cs.uwaterloo.ca/~jhoey/teaching/cs486/projects.html](https://cs.uwaterloo.ca/~jhoey/teaching/cs486/projects.html)

GenAI Policy

GenAI (ChatGPT, etc.) are known to be unreliable (<https://arxiv.org/pdf/2309.01219>), and their usage in this class is **strongly discouraged**.

Your ability to solve problems on a final exam will be severely impacted if you don't solve assignment problems yourself using the methods we learn in class. **Any usage of these tools must be declared on all submitted work. Usage of GenAI without declaring exactly what it was used for will result in a reduced grade for the work.** Further, for research projects, these tools will generate the uniform content across users, and projects using GenAI may also lose marks for a lack of originality.

Overview of the Course

Overview of the Course

Introduction, Agents and AI

Deterministic AI

- Representation and reasoning
 - States and searching
 - Features and constraints
- Supervised learning
 - Regression
- Deterministic planning
 - ≈ end of midterm material

Uncertainty in AI

- Representation and reasoning
 - Uncertainty and probability
 - Neural networks
- Learning
 - Bayesian learning
 - Unsupervised learning
 - Deep learning
- Planning with uncertainty
 - Markov decision processes
 - Reinforcement learning
- Topics on language models
 - N-Gram language models
 - Neural language models

Academic Integrity and Intellectual Property

See also official course outline on LEARN Property of UW:

- Lecture content, spoken and written (and any audio/video recording thereof)
- Lecture handouts, presentations, and other materials prepared for the course (e.g., PowerPoint slides)
- Questions or solution sets from various types of assessments (e.g., assignments, quizzes, tests, final exams)
- Work protected by copyright (e.g., any work authored by the instructor or TA or used by the instructor or TA with permission of the copyright owner)

Sharing intellectual property without the intellectual property owner's permission is a violation of intellectual property rights

What is Artificial Intelligence (AI)?

What is Artificial Intelligence (AI)?

The synthesis and analysis of computational agents that act
intelligently

What is Artificial Intelligence (AI)?

The synthesis and analysis of computational agents that act *intelligently*

Intelligence is a matter of **degree**

The aspects that go into an agent acting intelligently include

- What it does is appropriate for its **circumstances**, its **goals**, and its **perceptual and computational limitations**
- It takes into account the **short-term and long-term consequences** of its actions, including the effects on society and the environment
- It **learns** from experience
- It is **flexible** to changing environments and changing goals

Goals of Artificial Intelligence Research

Goals of Artificial Intelligence Research

Scientific: Understanding the principles that making intelligent behaviour possible in natural or artificial systems

- **Analysis** of natural and artificial agents
- **Testing hypotheses** about the construction of intelligent agents
- **Building and experimenting with** intelligent agents

Subareas of AI

Machine learning: Designing algorithms that learn from data

Deep learning: Machine learning, with specific focus on neural networks

Natural language processing and **computational linguistics:**
Understanding and generating human language

Computer vision: Interpreting and generating visual information

Subareas of AI

Machine learning: Designing algorithms that learn from data

Deep learning: Machine learning, with specific focus on neural networks

Natural language processing and **computational linguistics:**
Understanding and generating human language

Computer vision: Interpreting and generating visual information

Robotics: Building robots that can act autonomously

Subareas of AI

Machine learning: Designing algorithms that learn from data

Deep learning: Machine learning, with specific focus on neural networks

Natural language processing and **computational linguistics:**
Understanding and generating human language

Computer vision: Interpreting and generating visual information

Robotics: Building robots that can act autonomously

Cognitive science: Understanding how humans and animals think, learn, and act

History of AI: The Turing Test



Alan Turing: 1912-1954 was primarily known for codebreaking in WW11 (<https://www.imdb.com/title/tt2084970/>)

He also asked '*can machines think?*' and proposed the **imitation game** in 1950

An interrogator is connected to one person and one machine via a terminal, therefore can't see her counterparts. Her task is to find out which of the two candidates is the machine, and which is the human only by asking them questions. If the machine can "fool" the interrogator, it is intelligent.

Conversations with ELIZA - (Weizembaum 1966, Lisp)

ME: Hi

ELIZA: How do you do? What brings you to see me?

ME: I'm thinking of taking a course in artificial intelligence.

ELIZA: Why do you say that?

ME: I'm worried about it

ELIZA: Are you worried often?

ME: when I take AI courses

ELIZA: What happened when you take ai courses?

ME: I usually spend long nights doing assignments

ELIZA: Earlier you said you felt worried?

ME: about AI courses

ELIZA: Maybe your plans have something to do with this.

ME: no

ELIZA: Why not?

ME: because I think you are a computer

ELIZA: Is it because you think I am a computer that you came to me?

ME: no

ELIZA: Does it bother you that I ask that?

ME: yes

Loebner Prize: A Turing Test Competition

Fool all judges: \$1M prize

No one has ever won

Example interactions:

CONFEDERATE: Hello, my name is Sam and I am the human.

JUDGE: Hi Sam, tell me which is larger, a TV screen or Belgium.

CONFEDERATE: Belgium is a little bit larger

JUDGE: Wljs0 lwjejd0 wleje elwje0 wjeol, isn't it?

CONFEDERATE: Is that Polish?

JUDGE: Which is further away from this apartment, Queens or the Queen of England

CONFEDERATE: If the Queen is visiting Manhattan, she is closer.

Would ChatGPT win the Loebner Prize?

SHRDLU: A Program for Blocks World (Winograd, 1968)

First autonomous system supporting physical interaction

Blocks world only

Example conversation:

Pick up a big red block

OK

grasp the pyramid


I DON'T UNDERSTAND WHICH
PYRAMID YOU MEAN

What does the box contain?

THE BLUE PYRAMID AND THE
BLUE BLOCK.

What is the pyramid supported by?


THE BOX.



Game Playing

Two-player, zero-sum games

Game tree for tic-tac-toe:



Go

The game tree has more than 3^{361} nodes


- 3: Possible state for a single cell
- 361: Number of cells (19×19) on the board

Google DeepMind: AlphaGo (2016)

<https://www.nature.com/articles/nature16961>

March 2016: AlphaGo beats Lee Sedol in 4/5 games

May 2017: AlphaGo beats Ke Jie in 3/3 games



StarCraft

A multi-agent problem with

- Imperfect information (partially observed map)
- Large action space (10^8 possibilities)
- Delayed credit assignment



StarCraft

A multi-agent problem with


- Imperfect information (partially observed map)
- Large action space (10^8 possibilities)
- Delayed credit assignment

Google DeepMind:

<https://www.nature.com/articles/s41586-019-1724-z>



Robotics



Shakey
SRI
1970



José
UBC
2000



HRP-4C
AIST - Hiroshi Ishiguro
2010

Robocup

Soccer playing robots

Robocup 2023:

<https://www.youtube.com/watch?v=Ub1Z02dVKXM>



Autonomous Cars

“Stanley” (team led by Sebastian Thrun from Stanford) won the 2005 DARPA Grand Challenge



Autonomous Cars: 20 Years

Waymo-operated Jaguar
I-Pace in SF (2023)



Autopilot from Tesla



Are Self-Driving Cars (or AlphaGo etc.) Intelligent?

Recap: AI cares about the synthesis and analysis of computational agents that act *intelligently*

Intelligence is a matter of **degree**

The aspects that go into an agent acting intelligently include

- What it does is appropriate for its **circumstances**, its **goals**, and its **perceptual and computational limitations**
- It takes into account the **short-term and long-term consequences** of its actions, including the effects on society and the environment
- It **learns** from experience
- It is **flexible** to changing environments and changing goals

