



CS 886 Deep Learning for Biotechnology

Ming Li

CONTENT

- 01. Word2Vec
- 02. Attention / Transformers
- 03. Pretraining: GPT / BERT
- 04. Deep learning applications in proteomics
- 05. Student presentations begin
-
-




Attention and Transformers

LECTURE TWO

Plan: We trace back history to see how attention and transformers have emerged

1. Basic models, related to transduction models and attention.
2. Encoder-Decoder model, using recurrent networks such as LSTM.
3. Transformer models are general models sufficient for almost all biotech applications (graph models may be treated to be special cases too).
4. For example, DeepMind AlphaFold2 uses depends on a transformer architecture to train an end-to-end model.
5. The transformer model also makes it easy for large scale biological data (pre)training.


1. Fully connected network, feedforward network



To learn the weights on the edges

2. CNN

A CNN is a neural network with some convolutional layers (and some other layers). A convolutional layer has a number of filters that do convolutional operation.



Convolutional layer

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

Input

These are the network parameters to be learned.

1	-1	-1
-1	1	-1
-1	-1	1

Filter 1

-1	1	-1
-1	1	-1
-1	1	-1

Filter 2

⋮ ⋮

Each filter detects a small pattern (3 x 3).



Convolution Operation

stride=1

1	-1	-1
-1	1	-1
-1	-1	1

Filter 1

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

Dot
product



3

-1

Input



Convolution

stride=1

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0


Input

1	-1	-1
-1	1	-1
-1	-1	1

Filter 1


3	-1	-3	-1
-3	1	0	-3
-3	-3	0	1
3	-2	-2	-1

3. RNN




Parameters to be learned:
 U, V, W

Simple RNN vs LSTM




(a) RNN




(b) LSTM

Encoder-Decoder machine translation



Encoder-Decoder LSTM structure for chatting



Attention




Image caption generation using attention

z^0 is initial parameter, it is also learned





Image caption generation using attention





Image caption generation using attention





Image caption generation using attention




A woman is throwing a frisbee in a park.




A dog is standing on a hardwood floor.



A stop sign is on a road with a mountain in the background.



A little girl sitting on a bed with a teddy bear.



A group of people sitting on a boat in the water.




A giraffe standing in a forest with trees in the background.

Kelvin Xu, Jimmy Ba, Ryan Kiros, Kyunghyun Cho, Aaron Courville, Ruslan Salakhutdinov, Richard Zemel, Yoshua Bengio, “Show, Attend and Tell: Neural Image Caption Generation with Visual Attention”, ICML, 2015


More new ideas:

1. ULM-FiT, pre-training, transfer learning in NLP
2. Recurrent models require linear sequential computation, hard to parallelize. ELMo, bidirectional LSTM.
3. In order to reduce such sequential computation, several models based on CNN are introduced, such as ConvS2S and ByteNet. Dependency for ConvS2S needs linear depth, and ByteNet logarithmic.
4. The transformer is the first transduction model relying entirely on self-attention to compute the representations of its input and output without using RNN or CNN.

Transformer




An Encoder Block: same structure, different parameters



Encoder

Note: The ffnn is independent for each word.
Hence can be parallelized.



Self Attention

First we create three vectors
by multiplying input embedding
(1x512)

x_i with three matrices (64x512):

$$q_i = x_i W^Q$$

$$K_i = x_i W^K$$

$$V_i = x_i W^V$$



Input

Thinking

Machines




512 → 4

Embedding



 x_1  x_2 

Queries


64 → 3

 q_1  q_2 64
512  W^Q

Keys

 k_1  k_2  W^K

Values

 v_1  v_2  W^V

Self Attention

Now we need to calculate a score to determine how much focus to place on other Parts of the input.

Input

Embedding


Queries

Keys

Values

Score

Thinking

x_1 


q_1 


k_1 


v_1 


$q_1 \cdot k_1 = 112$

Machines

x_2 

q_2 

k_2 

v_2 

$q_1 \cdot k_2 = 96$

Self Attention

Formula

$$\text{softmax}\left(\frac{\begin{matrix} 64 \times 64 \\ Q \end{matrix} \times \begin{matrix} 64 \times 512 \\ K^T \end{matrix}}{\sqrt{d_k}}\right) \begin{matrix} 64 \times 512 \\ V \end{matrix} = \begin{matrix} Z \end{matrix}$$

$d_k=64$ is dimension of key vector

Input

Embedding

Queries

Keys

Values

Score

Divide by 8 ($\sqrt{d_k}$)

Softmax





Softmax

X
Value

Sum

Thinking

Machines




 x_1  x_2  q_1  q_2  k_1  k_2  v_1  v_2  $q_1 \cdot k_1 = 112$ $q_1 \cdot k_2 = 96$

14

12


0.88

0.12

 $\tilde{v}_1$  $\tilde{v}_2$  $z_1 = 0.88v_1 + 0.12v_2$ z_1  z_2 


Multiple heads

1. It expands the model's ability to focus on different positions.
2. It gives the attention layer multiple “representation subspaces”



The output
is
expecting
only a 2x4
matrix,
hence,



1) Concatenate all the attention heads



2) Multiply with a weight matrix W^O that was trained jointly with the model

X

3) The result would be the Z matrix that captures information from all the attention heads. We can send this forward to the FFNN




If you want some more intuition on attention: watch <https://www.youtube.com/watch?v=-9vVhYEXeyQ>

Representing the input order (positional encoding)


The transformer adds a vector to each input embedding. These vectors follow a specific pattern that the model learns, which helps it determine the position of each word, or the distance between different words in the sequence. The intuition here is that adding these values to the embeddings provides meaningful distances between the embedding vectors once they're projected into Q/K/V vectors and during dot-product attention.

More on positional encoding:
https://kazemnejad.com/blog/transformer_architecture_positional_encoding/



Add and Normalize

In order to regulate the computation, this is a normalization layer so that each feature (column) have the same average and deviation.



Layer Normalization (Hinton)


Layer normalization normalizes the inputs across the features.

Batch Normalization

batch			Same for all training examples	
			mean	std
1	3	6	3	3
2	2	2	2	0
0	1	5	3	3
4	6	1	4	3
5	2	3	3	2
1	0	1	1	1


Layer Normalization

batch			Same for all feature dimensions	
			mean	std
1	3	6	2	2
2	2	2	3	2
0	1	5	3	2
4	6	1		
5	2	3		
1	0	1		



The complete transformer

The encoder-decoder attention is just like self attention, except it uses K, V from the top of encoder output, and its own Q





Attention and Transformers

Note: In decoder, the input is “incomplete” when calculating self-attention.

The solution is to set future unknown values with “-inf” .



Attention and Transformers

Which word in our vocabulary
is associated with this index?

Decoder's

Output

Linear

Layer

Get the index of the cell
with the highest value
(**argmax**)

log_probs



am

5

logits



Softmax


Linear

Decoder stack output



Decoding time step: 1 2 3 4 5 6

OUTPUT | am



But what about
Self-attention?

Training and the Loss Function

Untrained Model Output



Correct and desired output



a


am

I

thanks

student

<eos>



We can use cross Entropy.

We can also optimize two words at a time: using BEAM search: keep a few alternatives for the first word.

Cross Entropy and KL (Kullback-Leibler) divergence


- **Entropy**: $E(P) = - \sum_i P(i) \log P(i)$ - expected prefix-free code length (also optimal)
- **Cross Entropy**: $C(P) = - \sum_i P(i) \log Q(i)$ – expected coding length using optimal code for Q
- **KL divergence**:
 $D_{KL}(P \parallel Q) = \sum_i P(i) \log[P(i)/Q(i)] = \sum_i P(i) [\log P(i) - \log Q(i)]$, extra bits to code using Q rather than P
- **JSD** $(P \parallel Q) = \frac{1}{2} D_{KL}(P \parallel M) + \frac{1}{2} D_{KL}(Q \parallel M)$, $M = \frac{1}{2} (P + Q)$, symmetric KL

* JSD = Jensen-Shannon Divergency

Transformer Results

Table 2: The Transformer achieves better BLEU scores than previous state-of-the-art models on the English-to-German and English-to-French newstest2014 tests at a fraction of the training cost.

Model	BLEU		Training Cost (FLOPs)	
	EN-DE	EN-FR	EN-DE	EN-FR
ByteNet [18]	23.75			
Deep-Att + PosUnk [39]		39.2		$1.0 \cdot 10^{20}$
GNMT + RL [38]	24.6	39.92	$2.3 \cdot 10^{19}$	$1.4 \cdot 10^{20}$
ConvS2S [9]	25.16	40.46	$9.6 \cdot 10^{18}$	$1.5 \cdot 10^{20}$
MoE [32]	26.03	40.56	$2.0 \cdot 10^{19}$	$1.2 \cdot 10^{20}$
Deep-Att + PosUnk Ensemble [39]		40.4		$8.0 \cdot 10^{20}$
GNMT + RL Ensemble [38]	26.30	41.16	$1.8 \cdot 10^{20}$	$1.1 \cdot 10^{21}$
ConvS2S Ensemble [9]	26.36	41.29	$7.7 \cdot 10^{19}$	$1.2 \cdot 10^{21}$
Transformer (base model)	27.3	38.1	$3.3 \cdot 10^{18}$	
Transformer (big)	28.4	41.8	$2.3 \cdot 10^{19}$	



Literature & Resources for Transformers

Vaswani et al. Attention is all you need. 2017.

Resources:

<https://nlp.seas.harvard.edu/2018/04/03/attention.html> (Excellent explanation of transformer model with codes.)

Jay Alammar, The illustrated transformer (from which I borrowed many pictures):

<http://jalammar.github.io/illustrated-transformer/>